

A. Feder Cooper, Ph.D.

Email: afedercooper@gmail.com

Website: afedercooper.info

I research a variety of topics in reliable, scalable machine learning. My contributions span privacy, security, and evaluation of generative-AI systems, MLSys, and uncertainty estimation. I also do work in tech policy and law, and spend a lot of time finding ways to effectively communicate the capabilities and limits of AI/ML to interdisciplinary audiences and the public. My research has received spotlights, orals, and best-paper accolades at top AI/ML and computing venues, including *NeurIPS*, *ICML*, *AAAI*, and *AIES*. Law collaborations on copyright and Generative AI have been lauded as “landmark” work among technology law scholars and the popular press.

Positions and Appointments

Co-lead, [The Center for Generative AI, Law, and Policy Research](#)

Assistant Professor, Department of Computer Science, Yale University (starting July 2026)

Postdoctoral Researcher, Microsoft Research

Postdoctoral Affiliate, Stanford HAI, RegLab, and CRFM

Faculty Associate at the Berkman Klein Center for Internet & Society at Harvard University

Education

Cornell University

Ph.D., Computer Science

2018 – 2024

[“Between Randomness and Arbitrariness: Some Lessons for Reliable Machine Learning at Scale”](#)

M.S., Computer Science

2021

Columbia University

B.A., Computer Science and Archaeology, Phi Beta Kappa, *summa cum laude*

2014

Papers

*Denotes co-first author or equal contribution; †lead author (not specified if same as first author(s))

A. Feder Cooper, Aaron Gokaslan, Amy B. Cyphert, Christopher De Sa, Mark A. Lemley, Daniel E. Ho, Percy Liang. “Extracting memorized pieces of (copyrighted) books from open-weight language models.” 2025. [\[link\]](#)

Jamie Hayes*, Ilia Shumailov, Christopher A. Choquette-Choo, Matthew Jagielski, George Kaissis, Katherine Lee, Milad Nasr, Sahra Ghalebikesabi, Niloofar Mireshghallah, Meenatchi Sundaram Mutu Selva Annamalai, Igor Shilov, Matthieu Meeus, Yves-Alexandre de Montjoye, Franziska Boenisch, Adam Dziedzic, and **A. Feder Cooper***. “Strong Membership Inference Attacks on Massive Datasets and (Moderately) Large Language Models.” 2025. [\[link\]](#)

Nikhil Kandpal*, Brian Lester*, Colin Raffel*, Sebastian Majstorovic, Stella Biderman, Baber Abbasi, Luca Soldaini, Enrico Shippole, **A. Feder Cooper**, Aviya Skowron, John Kirchenbauer, Shayne Longpre, Lintang Sutawika, Alon Albalak, Zhenlin Xu, Guilherme Penedo, Loubna Ben Allal, Elie Bakouch, John David Pressman, Honglu Fan, Dashiell Stander, Guangyu Song, Aaron Gokaslan, Tom Goldstein, Brian R. Bartoldson, Bhavya Kailkhura, and Tyler Murray. “The Common Pile v0.1: An 8TB Dataset of Public Domain and Openly Licensed Text.” 2025. [\[link\]](#)

Jamie Hayes*, Marika Swanberg, Harsh Chaudhari, Itay Yona, Ilia Shumailov, Milad Nasr, Christopher A. Choquette-Choo, Katherine Lee, and **A. Feder Cooper***. “Measuring memorization in language models via probabilistic extraction.” 2025 *Annual Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics (NAACL ’25)*. [\[link\]](#)

A. Feder Cooper* and James Grimmelmann*. “The Files are in the Computer: On Copyright, Memorization, and Generative AI.” Forthcoming, *Chicago-Kent Law Review*, 2025 (written in 2024). [\[link\]](#)

Non-archival versions: *2nd Workshop on Generative AI and Law (GenLaw)* at *The Forty-first International Conference on Machine Learning (ICML ’24)*; *Intellectual Property Law Scholars Conference 2024 (IPSC ’24)*; *The 4th ACM Symposium on Computer Science and Law (CSLAW ’25)*, 33% acceptance rate

A. Feder Cooper*, Katherine Lee*, and James Grimmelmann*. “Talkin’ Bout AI Generation: Copyright and the Generative-AI Supply Chain.” Forthcoming, *Journal of the Copyright Society*, 2025 (written in 2023). [[link](#)]

Preliminary version: **A. Feder Cooper***, Katherine Lee*, and James Grimmelmann*. “Talkin’ Bout AI Generation: Copyright and the Generative-AI Supply Chain (The Short Version).” *The 3rd ACM Symposium on Computer Science and Law (CSLAW ’24)*, 24% acceptance rate, **Long Presentation**. [[link](#)]

Preliminary version: Katherine Lee, **A. Feder Cooper**, and James Grimmelmann. “Beyond Generation: Privacy and Copyright Consequences of Retrieval and Attribution Models.” *Privacy Law Scholars Conference 2023 (PLSC ’23)*.

Milad Nasr*, Javier Rando*, Nicholas Carlini*, Jonathan Hayase, Matthew Jagielski, **A. Feder Cooper**, Daphne Ippolito, Christopher A. Choquette-Choo, Florian Tramèr, and Katherine Lee†. “Scalable Extraction of Training Data from Aligned, Production Language Models.” *The Thirteenth International Conference on Learning Representations (ICLR ’25)*, 32% acceptance rate. [[link](#)]

A. Feder Cooper*†, Christopher A. Choquette-Choo*, Miranda Bogen*, Matthew Jagielski*, Katja Filippova*, Ken Ziyu Liu*, Alexandra Chouldechova, Jamie Hayes, Yangsibo Huang, Niloofar Miresghallah, Ilia Shumailov, Eleni Triantafillou, Peter Kairouz, Nicole Mitchell, Percy Liang, Daniel E. Ho, Yejin Choi, Sanmi Koyejo, Fernando Delgado, James Grimmelmann, Vitaly Shmatikov, Christopher De Sa, Solon Barocas, Amy Cyphert, Mark Lemley, danah boyd, Jennifer Wortman Vaughan, Miles Brundage, David Bau, Seth Neel, Abigail Z. Jacobs, Andreas Terzis, Hanna Wallach, Nicolas Papernot, and Katherine Lee†. “Machine Unlearning Doesn’t Do What You Think: Lessons for Generative AI Policy, Research, and Practice.” 2024. [[link](#)]

Preliminary versions: *2nd Workshop on Generative AI and Law (GenLaw)* at *The Forty-first International Conference on Machine Learning (ICML ’24)*; *The 4th ACM Symposium on Computer Science and Law (CSLAW ’25)*, Works-in-Progress Track

Hanna Wallach*, Meera Desai, Nicholas J. Pangakis, **A. Feder Cooper**, Angelina Wang, Solon Barocas, Alexandra Chouldechova, Chad Atalla, Su Lin Blodgett, Emily Corvi, P. Alexander Dow, Jean Garcia-Gathright, Alexandra Olteanu, Stefanie Reed, Emily Shang, Dan Vann, Jenn Wortman Vaughan, Matthew Vogel, Hannah Washington, and Abigail Z. Jacobs†. “Evaluating Generative AI Systems is a Social Science Measurement Challenge.” *Workshop on Evaluating Evaluations (EvalEval)* at the *Conference on Neural Information Processing Systems 37 (NeurIPS ’24)*, **Oral**, < 10% of submitted papers. (Archival version under submission.) [[link](#)]

A. Feder Cooper. “Between Randomness and Arbitrariness: Some Lessons for Reliable Machine Learning at Scale.” Ph.D. Dissertation, Cornell University, August 2024. [[link](#)]

Blog version: **A. Feder Cooper**. “Between Randomness and Arbitrariness: Some Lessons for Reliable Machine Learning at Scale (The Short Version),” August 2024. [[link](#)]

Nicholas Carlini, Daniel Paleka, Krishnamurthy Dj Dvijotham, Thomas Steinke, Jonathan Hayase, **A. Feder Cooper**, Katherine Lee, Matthew Jagielski, Milad Nasr, Arthur Conmy, Eric Wallace, David Rolnick, and Florian Tramèr. “Stealing Part of a Production Language Model.” *International Conference on Machine Learning 2024 (ICML ’24)*, **Best Paper Award**, ~ 0.1% of submitted papers. [[link](#)]

Daniel McDuff, Tim Korjakow, Scott Cambo, Jesse Josua Benjamin, Jenny Lee, Yacine Jernite, Carlos Muñoz Ferrandis, Aaron Gokaslan, Alek Tarkowski, Joseph Lindley, **A. Feder Cooper**, and Danish Contractor. “On the Standardization of Behavioral Use Clauses and Their Adoption for Responsible Licensing of AI.” *International Conference on Machine Learning 2024 (ICML ’24)*, 27.5% acceptance rate. [[link](#)]

Aaron Gokaslan†, **A. Feder Cooper**, Jasmine Collins, Landan Seguin, Austin Jacobson, Mihir Patel, Jonathan Frankle, Cory Stephenson, and Volodymyr Kuleshov. “CommonCanvas: Open Diffusion Models Trained on Creative-Commons Images.” *Conference on Computer Vision and Pattern Recognition 2024 (CVPR ’24)*, 23.6% acceptance rate. [[link](#)]

Preliminary versions: *Workshop on ML for Creativity and Design*; *Workshop on Diffusion Models* at the *Conference on Neural Information Processing Systems 36 (NeurIPS ’23)*.

A. Feder Cooper†, Katherine Lee, Madiha Zahrah Choksi, Solon Barocas, Christopher De Sa, James Grimmelmann, Jon Kleinberg, Siddhartha Sen, and Baobao Zhang. “Arbitrariness and Social Prediction: The Confounding Role of Variance in Fair Classification.” *38th AAAI Conference on Artificial Intelligence (AAAI ’24)* (written in 2022), **Best Student Paper Honorable Mention**, < 1% of submitted papers. [[link](#)]

Preliminary version: *Workshop on Algorithmic Fairness through the Lens of Time* at the *Conference on Neural Information Processing Systems 36 (NeurIPS '23)*.

A. Feder Cooper[†], Wentao Guo*, Khiem Pham*, Tiancheng Yuan, Charlie F. Ruan, Yucheng Lu, and Christopher De Sa[†]. “Coordinating Distributed Example Orders for Provably Accelerated Training.” *Conference on Neural Information Processing Systems 36 (NeurIPS '23)*, 26% acceptance rate. [[link](#)]

Preliminary version: *3rd Workshop on Data-centric Machine Learning Research* at *The Fortieth International Conference on Machine Learning (ICML '23)*.

Kweku Kwegyir-Aggrey, **A. Feder Cooper**, Jessica Dai, John Dickerson, Keegan Hines, and Suresh Venkatasubramanian[†]. “Repairing Regressors for Fair Classification at Any Decision Threshold.” *Workshop on Optimal Transport and Machine Learning; Workshop on Algorithmic Fairness through the Lens of Time* at the *Conference on Neural Information Processing Systems 36 (NeurIPS '23)*, **Oral**. [[link](#)]

Milad Nasr*, Nicholas Carlini*, Jonathan Hayase, Matthew Jagielski, **A. Feder Cooper**, Daphne Ippolito, Christopher A. Choquette-Choo, Eric Wallace, Florian Tramèr, and Katherine Lee[†]. “Scalable Extraction of Training Data from (Production) Language Models.” 2023. [[link](#)]

A. Feder Cooper[†], Katherine Lee*, James Grimmelmänn*, Daphne Ippolito*, Christopher Callison-Burch, Christopher A. Choquette-Choo, Niloofar Mireshghallah, Miles Brundage, David Mimno, Madiha Zahrah Choksi, Jack M. Balkin, Nicholas Carlini, Christopher De Sa, Jonathan Frankle, Deep Ganguli, Bryant Gipson, Andres Guadamuz, Swee Leng Harris, Abigail Z. Jacobs, Elizabeth Joh, Gautam Kamath, Mark Lemley, Cass Matthews, Christine McLeavey, Corynne McSherry, Milad Nasr, Paul Ohm, Adam Roberts, Tom Rubin, Pamela Samuelson, Ludwig Schubert, Kristen Vaccaro, Luis Villa, Felix Wu, and Elana Zeide. “Report of the 1st Workshop on Generative AI and Law.” 2023. [[link](#)]

A. Feder Cooper*, Katherine Lee*, James Grimmelmänn, and Daphne Ippolito. “AI and Law: The Next Generation (An explainer series).” 2023. [[link](#)]

A. Feder Cooper, Madiha Zahrah Choksi, Katherine Lee, Solon Barocas, Christopher De Sa, James Grimmelmänn, Jon Kleinberg, Siddhartha Sen, and Baobao Zhang. “Distribution Justice: Variance, Uncertainty, and Rules in Machine Learning and Law.” *Data (Re)Makes the World Conference*, Information Society Project at Yale Law School, < 17% acceptance rate; *Privacy Law Scholars Conference*, 2023 (Archival version in progress). [[link](#)]

A. Feder Cooper, Jonathan Frankle, and Christopher De Sa. “Non-Determinism and the Lawlessness of Machine Learning Code.” *The 2nd ACM Symposium on Computer Science and Law (CSLAW '22)*, **Long Presentation**. [[link](#)]

A. Feder Cooper and Karen Levy. “Fast or Accurate? Governing Conflicting Goals in Highly Autonomous Vehicles.” *Colorado Technology Law Journal*, Vol. 20, 2022. [[link](#)]

A. Feder Cooper[†], Emanuel Moss*, Benjamin Laufer, and Helen Nissenbaum[†]. “Accountability in an Algorithmic Society: Relationality, Responsibility, and Robustness in Machine Learning.” *Proceedings of the 5th ACM Conference on Fairness, Accountability, and Transparency (FAccT '22)*, 25% acceptance rate. [[link](#)]

A. Feder Cooper and Gili Vidan. “Making the Unaccountable Internet: The Changing Meaning of Accounting in the Early ARPANET.” *Proceedings of the 5th ACM Conference on Fairness, Accountability, and Transparency (FAccT '22)*, 25% acceptance rate. [[link](#)]

Benjamin Laufer, **A. Feder Cooper**, Sameer Jain, Jon Kleinberg[†], and Hoda Heidari[†]. “Four Years of FAccT: A Reflexive, Mixed-Methods Analysis of Research Contributions, Shortcomings, and Future Prospects.” *Proceedings of the 5th ACM Conference on Fairness, Accountability, and Transparency (FAccT '22)*, 25% acceptance rate. [[link](#)]

A. Feder Cooper, Solon Barocas, Karen Levy, and Gili Vidan[†]. “‘We have met the enemy and it is us’: Debating the ethics of computing in the pages of CACM.” *2022 Workshop of the The Special Interest Group for Computing, Information, and Society (SIGCIS '22)*, **Talk**.

A. Feder Cooper, Yucheng Lu, Jessica Zosa Forde, and Christopher De Sa[†]. “Hyperparameter Optimization Is Deceiving Us, and How to Stop It.” *Conference on Neural Information Processing Systems 34 (NeurIPS '21)*, < 26% acceptance rate. [[link](#)]

Preliminary version: *Workshop on Robust Machine Learning* at *The Ninth International Conference on Learning Representations (ICLR '21)*.

A. Feder Cooper, Karen Levy, and Christopher De Sa. “Accuracy-Efficiency Trade-Offs and Accountability in Distributed ML Systems.” *Proceedings of the 2021 ACM Conference on Equity and Access in Algorithms, Mechanisms, and Optimization (EAAMO '21)*, **Oral**, < 10% of submitted papers. [\[link\]](#)

Preliminary version: **A. Feder Cooper**, Karen Levy, and Christopher De Sa. “Regulating Accuracy-Efficiency Trade-Offs in Distributed Machine Learning Systems.” *Workshop on Law and Machine Learning at The Thirty-seventh International Conference on Machine Learning (ICML '20)*, **Oral**.

Preliminary version: **A. Feder Cooper**. “Imperfection is the Norm: A Computer Systems Perspective on IoT and Enforcement.” *(Im)Perfect Enforcement Conference*, Information Society Project at Yale Law School, **Plenary session**, 2019. [\[link\]](#)

A. Feder Cooper, Maria Antoniak[†], Christopher De Sa, Marilyn Migiel[†], and David Mimno. “‘Tecnologica cosa’: Modeling Storyteller Personalities in Boccaccio’s *Decameron*.” *SIGHUM Workshop on Computational Linguistics for Cultural Heritage, Social Sciences, Humanities and Literature at The 2021 Conference on Empirical Methods in Natural Language Processing (EMNLP '21)*. [\[link\]](#)

A. Feder Cooper*, Jessica Zosa Forde*, Kweku Kwegyir-Aggrey, Christopher De Sa[†], and Michael Littman[†]. “Model Selection’s Disparate Impact in Real-World Deep Learning Applications.” *Workshop on the Science and Engineering of Deep Learning at The Ninth International Conference on Learning Representations (ICLR '21)*, **Oral**. [\[link\]](#)

A. Feder Cooper and Ellen Abrams. “Emergent Unfairness in Algorithmic Fairness-Accuracy Trade-Off Research.” *Proceedings of the 2021 AAAI/ACM Conference on Artificial Intelligence, Ethics, and Society (AIES '21)*, **Oral**, < 10% of submitted papers. [\[link\]](#)

A. Feder Cooper*, Ruqi Zhang*, and Christopher De Sa[†]. “Asymptotically Optimal Exact Minibatch Metropolis-Hastings.” *Conference on Neural Information Processing Systems 33 (NeurIPS '20)*, **Spotlight**, < 3% of submitted papers. [\[link\]](#)

Ruqi Zhang, **A. Feder Cooper**, and Christopher De Sa[†]. “AMAGOLD: Amortized Metropolis Adjustment for Efficient Stochastic Gradient MCMC.” *Proceedings of the Twenty-third International Conference on Artificial Intelligence and Statistics (AISTATS '20)*. [\[link\]](#)

Selected Honors and Awards

ICML Best Paper Award, [Stealing Part of a Production Model](#) (10 of 9473 (~0.1%) submitted papers) (2024)

AAAI Best Student Paper Honorable Mention, [Arbitrariness and Social Prediction: The Confounding Role of Variance in Fair Classification](#) (< 1% of submitted papers) (2024)

Oral Presentation, [Repairing Regressors for Fair Classification at Any Decision Threshold](#), *AFT Workshop at NeurIPS Workshop Oral Presentation* (4/25 accepted papers) (2023)

Meta Ph.D. Fellowship Finalist (2023) (short-listed out of 3200 candidates)

Outstanding *ICML* Reviewer (\leq 10% of reviewers) (2022)

Cornell Department of Computer Science Service Award (2022)

Rising Star in EECS, MIT (2021)

Digital Life Initiative Doctoral Fellowship (2021 – 2022)

AIES Oral Presentation, [Emergent Unfairness in Algorithmic Fairness-Accuracy Trade-Off Research](#) (< 10% of submitted papers) (2021)

Oral Presentation, [Model Selection’s Disparate Impact in Real-World Deep Learning Applications](#), *SEDL Workshop at ICLR* (2/15 accepted papers) (2021)

Fellowship Finalist, Two Sigma Ph.D. Diversity Fellowship (2021)

Fellowship Finalist, OpenPhil AI Ph.D. Fellowship (2021)

NeurIPS Spotlight, [Asymptotically Optimal Exact Minibatch Metropolis-Hastings](#) (< 3%) (2020)

Cornell University Fellowship (2018 – 2019)

Conference and Workshop Organization

4th ACM Symposium on Computer Science and Law (CS&Law '25). Executive committee member, publicity co-chair. [\[link\]](#)

2nd Workshop on Generative AI and Law (GenLaw '24). Held at *ICML 2024* in Vienna, Austria. Lead co-organizer with Katherine Lee. [\[link\]](#)

Evaluating Generative AI Systems: the Good, the Bad, and the Hype (GenLaw DC '24). Lead co-organizer. Sponsored by the K&L Gates initiative in Ethics and Computational Technologies at Carnegie Mellon University (CMU), and organized by the GenLaw Center (A. Feder Cooper, Katherine Lee, James Grimmelmann), Carnegie Mellon's K&L Gates Initiative (Hoda Heidari), the Georgetown Law Center (Paul Ohm), and the Center for Democracy and Technology (Alexandra Reeve Givens, Miranda Bogen). [\[link\]](#)

1st Workshop on Generative AI and Law (GenLaw '23). Held at *ICML 2023* in Honolulu, Hawai'i, 26% acceptance rate. Lead co-organizer with Katherine Lee; co-organized with Dr. Niloofar Miresghallah, Madiha Zahrah Choksi, Prof. James Grimmelmann, Prof. David Mimno, and Dr. Deep Ganguli [\[link\]](#)

Invited Talks, Panels, and Workshops

Invited speaker, [Workshop on Machine Unlearning for Generative AI](#), *International Conference on Machine Learning*, July 2025 (Vancouver, Canada).

Invited guest expert, "Machine Un-learning and Technical Guardrails," Future of Privacy Forum's Technologist Roundtable for Policymakers, May 28, 2025 (Virtual).

"Measuring memorization in language models via probabilistic extraction." Percy Liang's research group (p-lambda), Stanford University, March 3, 2025 (Virtual – Stanford, CA).

AI Governance Tutorial (with Katherine Lee, 28th Annual BTLJ-BCLT Spring Symposium: AI Governance at the Crossroads, February 27, 2025 (Berkeley, CA).

"Machine Unlearning Doesn't Do What You Think: Lessons for Generative AI Policy, Research, and Practice."

Guest Lecture for CPSC 310: Technology, Power, and Security: Political Challenges of the Computer Age, Yale University, April 8, 2025 (New Haven, CT).

Invited talk, NYU Law Innovation Policy Colloquium, April 3, 2025 (New York, NY).

Microsoft Research, January 13, 2025 (New York, NY).

Future of Privacy Forum, January 13, 2025 (Virtual – Washington, DC).

Dan Ho's research group (RegLab), January 8, 2025 (Virtual – Stanford, CA).

Meta AI Policy Roundtable, December 9, 2024 (Virtual – Menlo Park, CA).

Invited panelist, "Art & Copyright." *Workshop on Generative AI and Creativity* at *NeurIPS '24*, December 14, 2024 (Vancouver, BC, CA).

Invited panelist, "AI and Data Governance." *Queer in AI Workshop* at *NeurIPS '24*, December 11, 2024 (Vancouver, BC, CA).

"Memorization, Copyright, and the Generative-AI Supply Chain."

Large Language Models: Methods and Applications course guest lecture, Carnegie Mellon University, November 19, 2024 (Virtual – Pittsburgh, PA).

Percy Liang's research group (p-lambda), Stanford University, October 7, 2024 (Virtual – Stanford, CA).

"Accountability and Generative-AI Supply Chains." AI Law and Policy course guest lecture, George Washington University Law School, October 2, 2024 (Virtual – Washington, DC).

"The Files are in the Computer: On Copyright, Memorization, and Generative AI."

Information Law Institute, New York University School of Law, September 18, 2024 (Virtual – New York, NY).

24th Annual Intellectual Property Scholars Conference, UC Berkeley School of Law, August 8, 2024 (Berkeley, CA).

Symposium on AI Disrupting Law, Chicago-Kent Law School, April 26, 2024 (Virtual – Chicago, IL).

Invited participant, “Transform: Copyright.” Harvard Law School. April 22, 2024 (Cambridge, MA).

“Copyright, Machine Learning Research, and the Generative-AI Supply Chain.” Invited talk with Katherine Lee. [Cohere for AI](#). April 3, 2024 (virtual).

“Reliable Measurement for ML at Scale.” Job talk, presented at several universities and industrial labs throughout Spring 2024.

“Extractable Memorization and Its Relationship to Copyright.” Prof. Yejin Choi’s Lab Group Meeting, University of Washington, January 21, 2024 (Seattle, WA).

“Talkin’ ’Bout AI Generation: Copyright and the Generative-AI Supply Chain.”

Ideas Lunch Speaker Series, Yale Law School, February 29, 2024 (Virtual – New Haven, CT).

Federated Learning Talks, Google Research, December 7, 2023 (Seattle, WA).

Carnegie Mellon University, November 30, 2023 (Pittsburgh, PA).

Invited panelist, “Governance & Accountability for ML: Existing Tools, Ongoing Efforts, & Future Directions.” *Conference on Neural Information Processing (NeurIPS)*, co-convened by Profs. Hoda Heidari and Emily Black, December 2023 (New Orleans, LA).

Invited panelist, “The Legal and Ethical Implications of Massive Scraped Training Datasets.” Foundation and Language Model (FLAME) Seminar, Carnegie Mellon University, December 1, 2023 (Pittsburgh, PA).

“More-Reliable Measurement in Algorithmic Fairness and Hyperparameter Optimization.” Machine Learning Department Faculty Seminar, Carnegie Mellon University, November 28, 2023 (Pittsburgh, PA).

“Talkin’ ’Bout AI Generation: Copyright, Machine Learning Research, and the Generative-AI Supply Chain.” [AS-SET Center for AI-Enabled Systems](#), University of Pennsylvania, School of Engineering and Applied Sciences, November 1, 2023 (Philadelphia, PA).

Invited participant, *Operationalizing the Measure Function of the NIST AI Risk Management Framework*, [Center for Advancing Safety of Machine Intelligence](#) & [NIST-NSF Institute for Trustworthy AI in Law and Society](#), & the [Federation of American Scientists](#), October 16-17, 2023 (Washington, DC).

Invited participant, *Sociotechnical Approaches to Measurement and Validation for Safety in AI*, [Center for Advancing Safety of Machine Intelligence](#), July 18-19, 2023 (Evanston, IL).

“Is My Prediction Arbitrary? Measuring Self-Consistency in Fair Classification.” Machine Learning Collective (MLC), June 2023 (virtual).

“Can Governance be Reconciled with Uncertainty in Machine Learning?” Hosted by the Center for the Study of Complex Systems, and co-sponsored by the Michigan Institute for Data Science (MIDAS) and the Center for Ethics, Society, and Computing (ESC), University of Michigan. March 23, 2023 (Ann Arbor, MI). [[link](#)]

“Uncertainty, Reliability, and Accountability in ML.” Facebook AI Research (Meta), Paris, December 2022 (virtual).

“Toward More Reliable Hyperparameter Optimization.” [[slides](#)]

Robotics Group, Brown Department of Computer Science, May 2022 (virtual)

Michael Carbin’s Lab, MIT EECS, February 2022 (virtual)

UberAI, February 2022 (virtual)

Invited speaker, Professor Rangita de Silva de Alwis’s *Spring Policy Lab: AI and Implicit Bias*, University of Pennsylvania Carey School of Law. April 2021, 2023 (virtual).

Featured in Professor Charles Isbell’s *NeurIPS 2020* Keynote Address: “You Can’t Escape Hyperparameters and Latent Variables: Machine Learning as a Software Engineering Enterprise.” December 2020. [[link](#)]

Affiliations

Stanford HAI, RegLab, and CRFM, Palo Alto, CA, August 2024 – present

Microsoft Research, New York, NY, July 2024 – present

Berkman Klein Center for Internet & Society, Harvard University, Fall 2023 – present

The GenLaw Center, July 2023 – present

Cornell University, Machine Learning Group, Ithaca and New York, NY, Fall 2018 – 2024

Google Research, Federated Learning Team, Seattle, WA and Remote, Fall 2023 - Spring 2024

CTRL-ALT Lab, Cornell Tech, 2022 – 2023

Artificial Intelligence Policy and Practice (AIPP), Cornell, Ithaca, NY, Spring 2018 – 2023

Microsoft Research, Remote (New York, NY), Summer 2022

Teaching

Guest lecturer, CS4787: Principles of Large-Scale Machine Learning | Fall 2022

Machine learning on GPUs. ML Accelerators.

Deployment and low-latency inference. Real-time learning. Deep neural network compression and pruning.

TA, CS5150: Software Engineering | Cornell University, Spring 2019 and Spring 2020

TA, CS6787: Advanced Machine Learning Systems | Cornell University, Fall 2018

Mentorship

I frequently serve in both unofficial and official capacities as a mentor to more junior Ph.D., master's, and undergraduate students (typically, undergraduates applying to graduate school). Formal research mentees while at Cornell:

Yimeng Zeng (Ph.D. Student at University of Pennsylvania CS); Jianan Canal Li (Ph.D. Student at University of California, Berkeley EECS); Dean Alvarez (Ph.D. Student at UIUC CS); Wentao Guo (Ph.D. Student at Princeton CS), Charlie Ruan, TC Yuan, Gary Wei

Selected Service

Area Chair for *COLM* (2025)

Publicity co-chair (with Aloni Cohen) for *CSLAW* (2025)

Reviewer / Program Committee for *NeurIPS* (2020-2024), *ICML* (2021-2023, 2025), *ICLR* (2022-2023, 2025), *TMLR* (2022-present), *AAAI* (2025), *CSLAW* (2025), *SATML* (2025), *ACM CACM* (2025), *COLM* (2024), *FAccT* (2021-2023), *JMLR* (2022), *AIES* (2021), *EAAMO* (2021)