

# A. Feder Cooper, Ph.D.

Email: [a.feder.cooper@yale.edu](mailto:a.feder.cooper@yale.edu)      Website: [afedercooper.info](http://afedercooper.info)

I research a variety of topics in reliable, scalable machine learning. My contributions span privacy, security, and evaluation of generative-AI systems, MLSys, and uncertainty estimation. I also do work in tech policy and law, and spend a lot of time finding ways to effectively communicate the capabilities and limits of AI/ML to interdisciplinary audiences and the public. My research has received spotlights, orals, and best-paper accolades at top AI/ML and computing venues, including *NeurIPS*, *ICML*, *AAAI*, and *AIES*. Law collaborations on copyright and Generative AI have been lauded as “landmark” work among technology law scholars and the popular press.

## Positions and Appointments

Assistant Professor, Department of Computer Science, Yale University (starting July 2026)

Founding Research Scientist, AVERI

Postdoctoral Affiliate, Stanford HAI, RegLab, and CRFM

Faculty Associate at the Berkman Klein Center for Internet & Society at Harvard University

## Education

### Cornell University

Ph.D., Computer Science

2018 – 2024

“[Between Randomness and Arbitrariness: Some Lessons for Reliable Machine Learning at Scale](#)”

M.S., Computer Science

2021

### Columbia University

B.A., Computer Science and Archaeology, Phi Beta Kappa, *summa cum laude*

2014

## Papers

\*Denotes co-first author or equal contribution; †lead author (not specified if same as first author)

36. Miles Brundage, Noemi Dreksler, Aidan Homewood, Sean McGregor, Patricia Paskov, Conrad Stosz, Girish Sastry, **A. Feder Cooper**, George Balston, Steven Adler, Stephen Casper, Markus Anderljung, Grace Werner, Soren Mindermann, Vasilios Mavroudis, Ben Bucknall, Charlotte Stix, Jonas Freund, Lorenzo Pacchiardi, Jose Hernandez-Orallo, Matteo Pistillo, Michael Chen, Chris Painter, Dean W. Ball, Cullen O’Keefe, Gabriel Weil, Ben Harack, Graeme Finley, Ryan Hassan, Scott Emmons, Charles Foster, Anka Reuel, Bri Teece, Yoshua Bengio, Daniel Reti, Rishi Bommasani, Cristian Trout, Ali Shahin Shamsabadi, Rajiv Dattani, Adrian Weller, Robert Trager, Jaime Sevilla, Lauren Wagner, Lisa Soder, Ketan Ramakrishnan, Henry Papadatos, Malcolm Murray, and Ryan Tovcimak. ”Frontier AI Auditing: Toward Rigorous Third-Party Assessment of Safety and Security Practices at Leading AI Companies.” 2026. [[link](#)] [arXiv](#)
35. Ahmed Ahmed\*, **A. Feder Cooper\***, Sanmi Koyejo, and Percy Liang. “Extracting books from production language models.” 2026. [[link](#)] [arXiv](#)
34. **A. Feder Cooper**, Aaron Gokaslan, Ahmed Ahmed, Amy B. Cyphert, Christopher De Sa, Mark A. Lemley, Daniel E. Ho, Percy Liang. “Extracting memorized pieces of (copyrighted) books from open-weight language models.” *ICML 2025 Workshop on Reliable and Responsible Foundation Models (R2-FM ’25)*. [Workshop](#) [Oral](#) (≈3% of submitted papers) [[link](#)]
33. Jamie Hayes\*, Ilia Shumailov, Christopher A. Choquette-Choo, Matthew Jagielski, George Kaassis, Milad Nasr, Meenatchi Sundaram Mutu Selva Annamalai, Niloofar Miresghallah, Igor Shilov, Matthieu Meeus, Yves-Alexandre de Montjoye, Katherine Lee, Franziska Boenisch, Adam Dziedzic, and **A. Feder Cooper**\*. “Exploring the limits of strong membership inference attacks on large language models.” *Conference on Neural Information Processing Systems 38 (NeurIPS ’25)*. [Proceedings](#) (24.52% overall acceptance rate) [[link](#)]

32. Nikhil Kandpal\*, Brian Lester\*, Colin Raffel\*, Sebastian Majstorovic, Stella Biderman, Baber Abbasi, Luca Soldaini, Enrico Shippole, **A. Feder Cooper**, Aviya Skowron, John Kirchenbauer, Shayne Longpre, Lintang Sutawika, Alon Albala, Zhenlin Xu, Guilherme Penedo, Loubna Ben Allal, Elie Bakouch, John David Pressman, Honglu Fan, Dashiell Stander, Guangyu Song, Aaron Gokaslan, Tom Goldstein, Brian R. Bartoldson, Bhavya Kailkhura, and Tyler Murray. “The Common Pile v0.1: An 8TB Dataset of Public Domain and Openly Licensed Text.” *Conference on Neural Information Processing Systems 38 (NeurIPS '25), Datasets and Benchmarks*. **Proceedings** (24.91% overall acceptance rate) [[link](#)]
31. **A. Feder Cooper\***<sup>†</sup>, Christopher A. Choquette-Choo\*, Miranda Bogen\*, Kevin Klyman\*, Matthew Jagielinski\*, Katja Filippova\*, Ken Ziyu Liu\*, Alexandra Chouldechova, Jamie Hayes, Yangsibo Huang, Eleni Triantafillou, Peter Kairouz, Nicole Mitchell, Niloofar Mireshghallah, Abigail Z. Jacobs, James Grimmelmann, Vitaly Shmatikov, Christopher De Sa, Ilia Shumailov, Andreas Terzis, Solon Barocas, Jennifer Wortman Vaughan, danah boyd, Yejin Choi, Sanmi Koyejo, Fernando Delgado, Percy Liang, Daniel E. Ho, Pamela Samuelson, Miles Brundage, David Bau, Seth Neel, Hanna Wallach, Amy B. Cyphert, Mark A. Lemley, Nicolas Papernot, and Katherine Lee<sup>\*†</sup>. “Machine Unlearning Doesn’t Do What You Think: Lessons for Generative AI Policy and Research.” *Conference on Neural Information Processing Systems 38 (NeurIPS '25), Position Track*. **Proceedings** **Oral** (≈0.4% of submitted papers). [[link](#)]
- Preliminary versions: *2nd Workshop on Generative AI and Law (GenLaw)* at *The Forty-first International Conference on Machine Learning (ICML '24)*; *The 4th ACM Symposium on Computer Science and Law (CSLAW '25)*, Works-in-Progress Track **Workshop**
30. Alexandra Chouldechova\*, **A. Feder Cooper\***, Solon Barocas, Abhinav Palia, Dan Vann, and Hanna Wallach. “Comparison requires valid measurement: Rethinking attack success rate comparisons in AI red teaming.” *Conference on Neural Information Processing Systems 38 (NeurIPS '25), Position Track*. **Proceedings** (5.7% overall acceptance rate)
29. Hanna Wallach\*, Meera Desai, **A. Feder Cooper**, Angelina Wang, Solon Barocas, Alexandra Chouldechova, Chad Atalla, Su Lin Blodgett, Emily Corvi, P. Alexander Dow, Jean Garcia-Gathright, Alexandra Olteanu, Nicholas J. Pangakis, Stefanie Reed, Emily Shang, Dan Vann, Jenn Wortman Vaughan, Matthew Vogel, Hannah Washington, and Abigail Z. Jacobs<sup>†</sup>. “Position: Evaluating Generative AI Systems is a Social Science Measurement Challenge.” *International Conference on Machine Learning 2025 (ICML '25)*. [[link](#)] **Proceedings** (19.7% overall acceptance rate)
- Preliminary versions: *Workshop on Evaluating Evaluations (EvalEval)* at the *Conference on Neural Information Processing Systems 37 (NeurIPS '24)*. [[link](#)] **Workshop** **Oral** (<10% of submitted papers).
28. Jamie Hayes\*, Marika Swanberg, Harsh Chaudhari, Itay Yona, Ilia Shumailov, Milad Nasr, Christopher A. Choquette-Choo, Katherine Lee, and **A. Feder Cooper\***. “Measuring memorization in language models via probabilistic extraction.” *2025 Annual Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics (NAACL '25)*. [[link](#)] **Proceedings** (~22% overall acceptance rate)
27. **A. Feder Cooper\*** and James Grimmelmann\*. “The Files are in the Computer: On Copyright, Memorization, and Generative AI.” *Chicago-Kent Law Review*, Vol. 100, 2025. [[link](#)] **Journal** (written 2024)
- Non-archival versions: *2nd Workshop on Generative AI and Law (GenLaw)* at *The Forty-first International Conference on Machine Learning (ICML '24)* **Workshop**; *Intellectual Property Law Scholars Conference 2024 (IPSC '24)* **Workshop**; *The 4th ACM Symposium on Computer Science and Law (CSLAW '25)* (33% overall acceptance rate)
26. Katherine Lee\*, **A. Feder Cooper\***, and James Grimmelmann\*. “Talkin’ ‘Bout AI Generation: Copyright and the Generative-AI Supply Chain.” *Journal of the Copyright Society*, Volume 72, 2025. [[link](#)] **Journal** (written 2023)
- Preliminary version: Katherine Lee\*, **A. Feder Cooper\***, and James Grimmelmann\*. “Talkin’ ‘Bout AI Generation: Copyright and the Generative-AI Supply Chain (The Short Version).” *The 3rd ACM Symposium on Computer Science and Law (CSLAW '24)*. [[link](#)] **Proceedings** (24% overall acceptance rate)
- Preliminary version: Katherine Lee, **A. Feder Cooper**, and James Grimmelmann. “Beyond Generation: Privacy and Copyright Consequences of Retrieval and Attribution Models.” *Privacy Law Scholars Conference 2023 (PLSC '23)*. **Workshop**

25. Milad Nasr\*, Javier Rando\*, Nicholas Carlini\*, Jonathan Hayase, Matthew Jagielski, **A. Feder Cooper**, Daphne Ippolito, Christopher A. Choquette-Choo, Florian Tramèr, and Katherine Lee<sup>†</sup>. “Scalable Extraction of Training Data from Aligned, Production Language Models.” *The Thirteenth International Conference on Learning Representations (ICLR ’25)*. [link] Proceedings (32% overall acceptance rate)
24. **A. Feder Cooper**. “Between Randomness and Arbitrariness: Some Lessons for Reliable Machine Learning at Scale.” Cornell University, August 2024. [link] Ph.D. Thesis  
Blog version: **A. Feder Cooper**. “Between Randomness and Arbitrariness: Some Lessons for Reliable Machine Learning at Scale (The Short Version),” August 2024. [link] Blog
23. Nicholas Carlini, Daniel Paleka, Krishnamurthy Dj Dvijotham, Thomas Steinke, Jonathan Hayase, **A. Feder Cooper**, Katherine Lee, Matthew Jagielski, Milad Nasr, Arthur Conmy, Eric Wallace, David Rolnick, and Florian Tramèr. “Stealing Part of a Production Language Model.” *International Conference on Machine Learning 2024 (ICML ’24)*, Best Paper Award. [link] Proceedings Best Paper ( $\approx 1\%$  of submitted papers)
22. Daniel McDuff, Tim Korjakow, Scott Cambo, Jesse Josua Benjamin, Jenny Lee, Yacine Jernite, Carlos Muñoz Ferrandis, Aaron Gokaslan, Alek Tarkowski, Joseph Lindley, **A. Feder Cooper**, and Danish Contractor. “On the Standardization of Behavioral Use Clauses and Their Adoption for Responsible Licensing of AI.” *International Conference on Machine Learning 2024 (ICML ’24)*. [link] Proceedings (27.5% overall acceptance rate)
21. Aaron Gokaslan<sup>†</sup>, **A. Feder Cooper**, Jasmine Collins, Landan Seguin, Austin Jacobson, Mihir Patel, Jonathan Frankle, Cory Stephenson, and Volodymyr Kuleshov. “CommonCanvas: Open Diffusion Models Trained on Creative-Commons Images.” *Conference on Computer Vision and Pattern Recognition 2024 (CVPR ’24)*. [link] Proceedings (23.6% overall acceptance rate)  
Preliminary versions: *Workshop on ML for Creativity and Design*; *Workshop on Diffusion Models* at the *Conference on Neural Information Processing Systems 36 (NeurIPS ’23)*. Workshop
20. **A. Feder Cooper**<sup>†</sup>, Katherine Lee, Madiha Zahra Choksi, Solon Barocas, Christopher De Sa, James Grimmelmann, Jon Kleinberg, Siddhartha Sen, and Baobao Zhang. “Arbitrariness and Social Prediction: The Confounding Role of Variance in Fair Classification.” *38th AAAI Conference on Artificial Intelligence (AAAI ’24)*. [link] Proceedings Best Student Paper Honorable Mention (<1% of submitted papers; written 2023)  
Preliminary version: *Workshop on Algorithmic Fairness through the Lens of Time* at the *Conference on Neural Information Processing Systems 36 (NeurIPS ’23)*. Workshop
19. **A. Feder Cooper**\*<sup>†</sup>, Wentao Guo\*, Khiem Pham\*, Tiancheng Yuan, Charlie F. Ruan, Yucheng Lu, and Christopher De Sa<sup>†</sup>. “Coordinating Distributed Example Orders for Provably Accelerated Training.” *Conference on Neural Information Processing Systems 36 (NeurIPS ’23)*. [link] Proceedings (26% overall acceptance rate)  
Preliminary version: *3rd Workshop on Data-centric Machine Learning Research* at *The Fortieth International Conference on Machine Learning (ICML ’23)*. Workshop
18. Kweku Kwegyir-Aggrey, **A. Feder Cooper**, Jessica Dai, John Dickerson, Keegan Hines, and Suresh Venkatasubramanian<sup>†</sup>. “Repairing Regressors for Fair Classification at Any Decision Threshold.” *Workshop on Optimal Transport and Machine Learning*; *Workshop on Algorithmic Fairness through the Lens of Time* at the *Conference on Neural Information Processing Systems 36 (NeurIPS ’23)*. [link] Workshop Oral
17. Milad Nasr\*, Nicholas Carlini\*, Jonathan Hayase, Matthew Jagielski, **A. Feder Cooper**, Daphne Ippolito, Christopher A. Choquette-Choo, Eric Wallace, Florian Tramèr, and Katherine Lee<sup>†</sup>. “Scalable Extraction of Training Data from (Production) Language Models.” 2023. [link] arXiv
16. **A. Feder Cooper**\*<sup>†</sup>, Katherine Lee\*, James Grimmelmann\*, Daphne Ippolito\*, Christopher Callison-Burch, Christopher A. Choquette-Choo, Niloofar Mireshghallah, Miles Brundage, David Mimno, Madiha Zahra Choksi, Jack M. Balkin, Nicholas Carlini, Christopher De Sa, Jonathan Frankle, Deep Ganguli, Bryant Gipson, Andres Guadamuz, Swee Leng Harris, Abigail Z. Jacobs, Elizabeth Joh, Gautam Kamath,

Mark Lemley, Cass Matthews, Christine McLeavey, Corynne McSherry, Milad Nasr, Paul Ohm, Adam Roberts, Tom Rubin, Pamela Samuelson, Ludwig Schubert, Kristen Vaccaro, Luis Villa, Felix Wu, and Elana Zeide. "Report of the 1st Workshop on Generative AI and Law." 2023. [link] arXiv

15. Katherine Lee\*, A. Feder Cooper\*, James Grimmelmann, and Daphne Ippolito. "AI and Law: The Next Generation (An explainer series)." 2023. [link] Blog
14. A. Feder Cooper, Madiha Zahrah Choksi, Katherine Lee, Solon Barocas, Christopher De Sa, James Grimmelmann, Jon Kleinberg, Siddhartha Sen, and Baobao Zhang. "Distribution Justice: Variance, Uncertainty, and Rules in Machine Learning and Law." *Data (Re)Makes the World Conference*, Information Society Project at Yale Law School (< 17% overall acceptance rate); *Privacy Law Scholars Conference*, 2023. [link] Workshop
13. A. Feder Cooper, Jonathan Frankle, and Christopher De Sa. "Non-Determinism and the Lawlessness of Machine Learning Code." *The 2nd ACM Symposium on Computer Science and Law (CSLAW '22)*. [link] Proceedings
12. A. Feder Cooper and Karen Levy. "Fast or Accurate? Governing Conflicting Goals in Highly Autonomous Vehicles." *Colorado Technology Law Journal*, Volume 20, 2022. [link] Journal
11. A. Feder Cooper\*†, Emanuel Moss\*, Benjamin Laufer, and Helen Nissenbaum†. "Accountability in an Algorithmic Society: Relationality, Responsibility, and Robustness in Machine Learning." *Proceedings of the 5th ACM Conference on Fairness, Accountability, and Transparency (FAccT '22)*. [link] Proceedings (25% overall acceptance rate)
10. A. Feder Cooper and Gili Vidan. "Making the Unaccountable Internet: The Changing Meaning of Accounting in the Early ARPANET." *Proceedings of the 5th ACM Conference on Fairness, Accountability, and Transparency (FAccT '22)*. [link] Proceedings (25% overall acceptance rate)
9. Benjamin Laufer, Sameer Jain\*, A. Feder Cooper\*, Jon Kleinberg†, and Hoda Heidari†. "Four Years of FAccT: A Reflexive, Mixed-Methods Analysis of Research Contributions, Shortcomings, and Future Prospects." *Proceedings of the 5th ACM Conference on Fairness, Accountability, and Transparency (FAccT '22)*. [link] Proceedings (25% overall acceptance rate)
8. A. Feder Cooper, Solon Barocas, Karen Levy, and Gili Vidan†. "'We have met the enemy and it is us': Debating the ethics of computing in the pages of CACM." *2022 Workshop of the The Special Interest Group for Computing, Information, and Society (SIGCIS '22)*. [link] Workshop
7. A. Feder Cooper, Yucheng Lu, Jessica Zosa Forde, and Christopher De Sa†. "Hyperparameter Optimization Is Deceiving Us, and How to Stop It." *Conference on Neural Information Processing Systems 34 (NeurIPS '21)*. [link] Proceedings (<26% overall acceptance rate)  
Preliminary version: *Workshop on Robust Machine Learning at The Ninth International Conference on Learning Representations (ICLR '21)*. Workshop
6. A. Feder Cooper, Karen Levy, and Christopher De Sa. "Accuracy-Efficiency Trade-Offs and Accountability in Distributed ML Systems." *Proceedings of the 2021 ACM Conference on Equity and Access in Algorithms, Mechanisms, and Optimization (EAAMO '21)*. [link] Proceedings Oral (<10% of submitted papers)  
Preliminary version: A. Feder Cooper, Karen Levy, and Christopher De Sa. "Regulating Accuracy-Efficiency Trade-Offs in Distributed Machine Learning Systems." *Workshop on Law and Machine Learning at The Thirty-seventh International Conference on Machine Learning (ICML '20)*. Oral  
Preliminary version: A. Feder Cooper. "Imperfection is the Norm: A Computer Systems Perspective on IoT and Enforcement." *(Im)Perfect Enforcement Conference*, Information Society Project at Yale Law School, 2019. [link] Workshop Plenary Session
5. A. Feder Cooper, Maria Antoniak†, Christopher De Sa, Marilyn Migiel†, and David Mimno. "'Tecnologica cosa': Modeling Storyteller Personalities in Boccaccio's *Decameron*." *SIGHUM Workshop on Computational Linguistics for Cultural Heritage, Social Sciences, Humanities and Literature at The 2021 Conference on Empirical Methods in Natural Language Processing (EMNLP '21)*. [link] Workshop

4. Jessica Zosa Forde\*, **A. Feder Cooper**\*, Kweku Kwegyir-Aggrey, Christopher De Sa†, and Michael Littman†. “Model Selection’s Disparate Impact in Real-World Deep Learning Applications.” *Workshop on the Science and Engineering of Deep Learning at The Ninth International Conference on Learning Representations (ICLR ’21)*. [[link](#)] **Workshop** **Oral**
3. **A. Feder Cooper** and Ellen Abrams. “Emergent Unfairness in Algorithmic Fairness-Accuracy Trade-Off Research.” *Proceedings of the 2021 AAAI/ACM Conference on Artificial Intelligence, Ethics, and Society (AIES ’21)*. [[link](#)] **Proceedings** **Oral** (< 10% of submitted papers)
2. Ruqi Zhang\*, **A. Feder Cooper**\*, and Christopher De Sa†. “Asymptotically Optimal Exact Minibatch Metropolis-Hastings.” *Conference on Neural Information Processing Systems 33 (NeurIPS ’20)*. [[link](#)] **Proceedings** **Spotlight** (< 3% of submitted papers)
1. Ruqi Zhang, **A. Feder Cooper**, and Christopher De Sa†. “AMAGOLD: Amortized Metropolis Adjustment for Efficient Stochastic Gradient MCMC.” *Proceedings of the Twenty-third International Conference on Artificial Intelligence and Statistics (AISTATS ’20)*. [[link](#)] **Proceedings**

## Selected Honors and Awards

*NeurIPS* Oral, **Machine Unlearning Doesn’t Do What You Think: Lessons for Generative AI Policy and Research** ( $\approx 0.4\%$  of submitted papers) (2025)

*ICML* Workshop Oral (R2-FM ’25), **Extracting memorized pieces of (copyrighted) books from open-weight language models** (6 of 175 ( $\approx 3\%$ ) submitted papers) (2025)

*ICML* Best Paper Award, **Stealing Part of a Production Model** (10 of 9473 ( $\approx 0.1\%$ ) submitted papers) (2024)

*AAAI* Best Student Paper Honorable Mention, **Arbitrariness and Social Prediction: The Confounding Role of Variance in Fair Classification** (< 1% of submitted papers) (2024)

Meta Ph.D. Fellowship Finalist (2023) (short-listed out of 3200 candidates)

Outstanding *ICML* Reviewer ( $\leq 10\%$  of reviewers) (2022)

Cornell Department of Computer Science Service Award (2022)

Rising Star in EECS, MIT (2021)

*AIES* Oral, **Emergent Unfairness in Algorithmic Fairness-Accuracy Trade-Off Research** (< 10% of submitted papers) (2021)

Fellowship Finalist, Two Sigma Ph.D. Diversity Fellowship (2021)

Fellowship Finalist, OpenPhil AI Ph.D. Fellowship (2021)

*NeurIPS* Spotlight, **Asymptotically Optimal Exact Minibatch Metropolis-Hastings** (< 3%) (2020)

Cornell University Fellowship (2018 – 2019)

## Conference, Workshop, and Tutorial Organization

**5th ACM Symposium on Computer Science and Law (CS&Law ’26)**. Executive committee member, publicity chair. Held at UC Berkeley (Berkeley, CA). [[link](#)]

**Data Privacy, Memorization, & Legal Implications in Generative AI: A Practical Guide.** *NeurIPS* ’25 Tutorial, December 2, 2025 (San Diego, CA).

**4th ACM Symposium on Computer Science and Law (CS&Law ’25)**. Executive committee member, publicity co-chair. Held at TU Munich (Munich, Germany). [[link](#)]

**2nd Workshop on Generative AI and Law (GenLaw ’24)**. Held at *ICML* 2024 in Vienna, Austria. Lead co-organizer with Katherine Lee. [[link](#)]

**Evaluating Generative AI Systems: the Good, the Bad, and the Hype (GenLaw DC '24).** Lead co-organizer. Sponsored by the K&L Gates initiative in Ethics and Computational Technologies at Carnegie Mellon University (CMU), and organized by the GenLaw Center (A. Feder Cooper, Katherine Lee, James Grimmelmann), Carnegie Mellon's K&L Gates Initiative (Hoda Heidari), the Georgetown Law Center (Paul Ohm), and the Center for Democracy and Technology (Alexandra Reeve Givens, Miranda Bogen). [[link](#)]

**1st Workshop on Generative AI and Law (GenLaw '23).** Held at *ICML 2023* in Honolulu, Hawai'i, 26% acceptance rate. Lead co-organizer with Katherine Lee; co-organized with Dr. Niloofar Miresghallah, Madiha Zahrah Choksi, Prof. James Grimmelmann, Prof. David Mimno, and Dr. Deep Ganguli [[link](#)]

## Recent invited Talks, Panels, and Workshops

“Extracting memorized pieces of (copyrighted) books from open-weight language models.” Guest Lecture, NYU, November 13, 2025 (Virtual – New York, NY).

Invited speaker, Workshop & Lecture Series on the Law & Economics of Innovation, ETH Zürich, October 28–29, 2025 (Zürich, Switzerland).

Invited participant, Symposium on AI and Law, Harvard Law School, October 24, 2025 (Cambridge, MA).

Invited keynote fireside chat, *5o Congreso Internacional “Nuevas fronteras de la Propiedad Intelectual y Tecnologías Disruptivas”*, Facultad de Derecho Universidad Austral, September 10, 2025 (Virtual – Pilar, Argentina)

“What Copyright Can Learn from Memorization Measurements of Language Models.”

Invited speaker, Information Technology and Innovation Foundation, August 2025 (Virtual).

Invited speaker, [Workshop on the Impact of Memorization on Trustworthy Foundation Models, International Conference on Machine Learning](#), July 2025 (Vancouver, Canada).

Invited panelist, [Workshop on Machine Unlearning for Generative AI, International Conference on Machine Learning](#), July 2025 (Vancouver, Canada).

Invited guest expert, “Machine Un-learning and Technical Guardrails,” Future of Privacy Forum’s Technologist Roundtable for Policymakers, July 2025 (Virtual).

“Measuring memorization in language models via probabilistic extraction.” Percy Liang’s research group (p-lambda), Stanford University, March 3, 2025 (Virtual – Stanford, CA).

AI Governance Tutorial (with Katherine Lee, 28th Annual BTLJ-BCLT Spring Symposium: AI Governance at the Crossroads, February 27, 2025 (Berkeley, CA).

“Machine Unlearning Doesn’t Do What You Think: Lessons for Generative AI Policy, Research, and Practice.”

Guest Lecture for CPSC 310: Technology, Power, and Security: Political Challenges of the Computer Age, Yale University, April 8, 2025 (New Haven, CT).

Invited talk, NYU Law Innovation Policy Colloquium, April 3, 2025 (New York, NY).

Microsoft Research, January 13, 2025 (New York, NY).

Future of Privacy Forum, January 13, 2025 (Virtual – Washington, DC).

Dan Ho’s research group (RegLab), January 8, 2025 (Virtual – Stanford, CA).

Invited panelist, “Art & Copyright.” *Workshop on Generative AI and Creativity* at *NeurIPS ’24*, December 14, 2024 (Vancouver, BC, CA).

## Affiliations

Co-lead, **The GenLaw Center**, July 2023–present

Faculty Associate (previously Affiliate Member), **Berkman Klein Center for Internet & Society**, Harvard University, Fall 2023 – present

Postdoctoral researcher, **Microsoft Research**, New York, NY, July 2024–December 2025

Lab member, **Cornell University, Machine Learning Group**, Ithaca, NY, Fall 2018–2024

Student researcher, **Google Research**, Seattle, WA, Fall 2023–Spring 2024

Group member, **CTRL-ALT Lab**, Cornell Tech, 2022–2023

Group member, **Artificial Intelligence Policy and Practice (AIPP)**, **Cornell**, Ithaca, NY, 2018–2023

Intern, **Microsoft Research**, remote, Summer 2022

## Teaching

Guest lecturer, CS4787: Principles of Large-Scale Machine Learning | Fall 2022

Machine learning on GPUs. ML Accelerators.

Deployment and low-latency inference. Real-time learning. Deep neural network compression and pruning.

TA, CS5150: Software Engineering | Cornell University, Spring 2019 and Spring 2020

TA, CS6787: Advanced Machine Learning Systems | Cornell University, Fall 2018

## Selected Service

Area Chair for *COLM* (2025, 2026), *ICML* (2026)

Publicity co-chair (with Aloni Cohen) for *CSLAW* (2025); publicity chair for *CSLAW* (2026)

Reviewer / Program Committee for *NeurIPS* (2020–2024), *ICML* (2021–2023, 2025), *ICLR* (2022–2023, 2025–2026), *TMLR* (2022–present), *AAAI* (2025), *CSLAW* (2025–2026), *SATML* (2025), *ACM CACM* (2025), *COLM* (2024), *FAccT* (2021–2023), *JMLR* (2022), *AIES* (2021), *EAAMO* (2021)