

What AI Benchmarks Actually Measure: Adapting Convergent and Discriminant Validity to Interrogate Fifty-Six AI Benchmarks

Meera Desai¹, Sang T. Truong², Hanna Wallach⁴, Alex Chouldechova⁵,
A. Feder Cooper^{2,3}, Jean Garcia-Gathright⁴, Daniel E. Ho², Abigail Z. Jacobs¹,
Sanmi Koyejo², Nicholas Pangakis⁴, Angelina Wang⁶

¹University of Michigan, ²Stanford University, ³Yale University,

⁴Microsoft Research, ⁵Abridge, ⁶Cornell Tech

Abstract

Benchmarks play a central role in the development and governance of models, yet it is often unclear whether they actually measure the concepts they purport to measure (e.g., reasoning, refusal). We adapt the lenses of convergent and discriminant validity from the social sciences into an approach for interrogating AI benchmarks, which we use to interrogate 56 capability and safety benchmarks using 53 models. We assign benchmarks with substantively similar purported concepts to a shared *assigned concept*, and ask whether model rankings on benchmarks with the same assigned concept correlate strongly with one another, and whether model rankings on benchmarks with different assigned concepts correlate less strongly. We ask analogous questions using model scores at the item-level, drawing on item response theory (IRT) models. We find that correlations between model rankings on benchmarks with the same assigned safety concepts are often weak, suggesting that these assigned concepts may be conceptualized inconsistently across benchmarks. For benchmarks with assigned capability concepts (e.g., reasoning, knowledge), model rankings are often as strongly correlated among benchmarks with the same assigned concept as between benchmarks with different assigned concepts, suggesting that different assigned capability concepts may not discriminate that well from one another. In some cases, benchmarks that share benchmark design elements (e.g., task structure, score format) correlate more strongly with one another than benchmarks with the same assigned concept. Finally, model rankings on some individual benchmarks correlate more strongly with model rankings on benchmarks with a different assigned concept than with model rankings on benchmarks sharing their own assigned concept, suggesting that these benchmarks may measure a different concept than they purport to. For example, model rankings on BBQ-accuracy correlate more strongly with model rankings on benchmarks labeled with reasoning than with model rankings on benchmarks that share its assigned concept, bias. To support future empirical work on the validity of benchmarks, we release our extensive dataset of model outputs and scores at the item- and benchmark-level.

1 Introduction

AI benchmarks are critical to the development and governance of models. Model scores on benchmarks are often treated as compact quantitative summaries of model capabilities (e.g., reasoning, summarization) and safety (e.g., offensive generation, bias), shaping consequential decisions around adoption, research priorities, and broader narratives about AI progress. However, a growing body of work has raised concerns about the *construct validity* of benchmarks—that is, whether they measure the concepts they purport to (Wallach et al., 2025; Alaa et al., 2025; Bean et al., 2025; Subramonian et al., 2023; Blodgett et al., 2021).

We adapt *convergent and discriminant validity*—two lenses used in the social sciences for assessing construct validity—into an approach for investigating these concerns. Traditionally, convergent validity asks whether measurements from a newly proposed measurement instrument (e.g., a benchmark) correlate strongly with measurements from an already validated instrument designed to measure the same or a similar concept, while discriminant validity asks whether those measurements correlate less strongly with measurements from instruments designed to measure dissimilar concepts (Campbell & Fiske, 1959; Adcock & Collier, 2001; Messick, 1989).

Benchmarks pose a few challenges to this setup. First, the concepts that benchmarks purport to measure are often underspecified (Raji et al., 2021). Because of this, low correlation between two reasoning benchmarks could indicate that one or both benchmarks are invalid, or it could indicate that the benchmarks conceptualize reasoning in distinct but valid ways. Second, few, if any, benchmarks have been validated using construct validity, so the convergent validity of a newly proposed benchmark can only be assessed through comparison with a benchmark that could itself be invalid.

Rather than assessing the convergent or discriminant validity of an individual benchmark against another possibly invalid benchmark, we instead assess the *convergence* among all benchmarks that purport to measure similar concepts and the *discrimination* between all benchmarks that purport to measure one group of similar concepts and all benchmarks that purport to measure a second group of similar concepts. Because we consider this more expansive set of comparisons, we use the terms convergence and discrimination to refer to our approach, rather than convergent and discriminant validity. We assign benchmarks with substantively similar purported concepts to a shared *assigned concept*. Where we do interrogate individual benchmarks, we compare model rankings on an individual benchmark to model rankings on all the benchmarks with the same (or different) assigned concept. This avoids relying on a single comparison, in the absence of a validated standard, and instead utilizes the set of benchmarks with the same assigned concept as a comparison point. We support our quantitative findings with qualitative follow-up, which demonstrates the utility of our approach for assessing individual benchmarks despite the differences in our setting.

We use our approach to analyze our dataset of model outputs and scores from 53 models on 56 capability and safety benchmarks. Our analyses yield four main findings. First, correlations between model rankings on benchmarks with the same assigned concept are often weak for safety benchmarks. This may reflect the multi-dimensional nature of these concepts, and supports other calls for more careful conceptualization and documentation (Blodgett et al., 2021; Wallach et al., 2025). Second, model rankings among benchmarks with assigned capability concepts (e.g., knowledge, reasoning) correlate as strongly with model rankings on benchmarks with different assigned capability concepts as with those on benchmarks that share an assigned capability concept, which suggests that these benchmarks may not be measuring discriminable concepts. Third, in some cases benchmarks with shared design elements (e.g., score format, task structure) correlate more strongly with one another than with benchmarks that share an assigned concept. Finally, some individual benchmarks may measure a different concept than the one they purport to measure—for example, model rankings on BBQ-accuracy, which purports to measure bias, correlates more strongly with model rankings on benchmarks assigned with reasoning than those from other benchmarks assigned with bias, suggesting it may measure reasoning rather than bias. Because BBQ-accuracy is often the only bias benchmark reported in recent commercial model releases (Table E.1), this leaves a significant gap for model evaluations.

Taken together, these findings illustrate the utility of convergent and discriminant validity as lenses for interrogating benchmarks, and we encourage their broader adoption in benchmark development. We make the following contributions:

- Approach for interrogating AI benchmarks inspired by convergent and discriminant validity.
- Dataset of model outputs and scores at the item- and benchmark-level across 56 benchmarks and 53 models, which we release on HuggingFace.¹

¹: <https://huggingface.co/datasets/madesai/what-ai-benchmarks-actually-measure>

- Empirical findings from our approach and dataset, revealing insights into the concepts that benchmarks purport to measure.

2 Related work

Validity issues affecting benchmarks are a longstanding and well-documented concern (e.g., [Blodgett et al. \(2021\)](#); [Reuel et al. \(2024\)](#); [Bean et al. \(2025\)](#)), raising the question of how well benchmarks measure the capability or safety concepts they purport to. In the social sciences, construct validity ([Adcock & Collier, 2001](#); [Messick, 1989](#)) offers a structured framework for addressing such validity concerns. Construct validity can be assessed through multiple lenses; we focus on two in this work: convergent validity and discriminant validity. Researchers have argued that convergent and discriminant validity are relevant to designing and interrogating AI benchmarks, and have begun to adapt them as such ([Wallach et al., 2025](#); [Alaa et al., 2025](#); [Salaudeen et al., 2025](#); [Bean et al., 2025](#); [Xiao et al., 2023](#); [Liu et al., 2024](#)).

Several recent efforts examine the validity of benchmarks empirically by analyzing patterns of agreement across benchmarks and models, though not explicitly through the lenses of convergent and discriminant validity. [Ren et al. \(2024\)](#) find that model rankings on many safety benchmarks correlate strongly with a general capability score for that model. We address a related but much broader question, examining benchmark correlation structure among benchmarks that purport to measure many concepts within the categories of capability and safety. We also assess the relationship between benchmarks using item-level model scores, drawing on IRT models. [Burnell et al. \(2023a\)](#) explore the latent structure of model performance via factor analysis, noting that BBQ ([Parrish et al., 2022](#)) loads more strongly on a reasoning factor than a bias factor. However, their focus is on the latent structure of model ability rather than construct validity or what benchmarks measure. [Qian et al. \(2026\)](#) analyze the model rankings produced by capability benchmarks purporting to measure similar concepts in order to study item quality within benchmarks. We extend the analysis to safety benchmarks and examine whether benchmarks measure the concepts they purport to.

We use item response theory (IRT) in this work, a method from psychometrics. In AI research, IRT is primarily used to characterize model ability ([Martínez-Plumed et al., 2019](#); [Liu et al., 2025](#)) and construct more efficient benchmarks through item selection ([Hempstead et al., 2004](#)) or difficulty analysis ([Martínez-Plumed et al., 2022](#)). We similarly draw on IRT but use it diagnostically to ask whether jointly estimating a latent trait on two benchmarks with the same assigned concept improves the latent trait’s prediction on item-level model scores.

3 Methodology

3.1 Dataset collection

Selecting benchmarks and models We use the term benchmark to describe a dataset paired with a scoring metric, following [Raji et al. \(2021\)](#). We collected a dataset of model outputs and scores on 56 benchmarks from 53 models. To collect our set of benchmarks, we used SafetyPrompts ([Röttger et al., 2025](#)), a recent large-scale review of safety benchmarks, and HELM ([Liang et al., 2022](#)), a large-scale benchmark study as initial sources. We then expanded to maximize coverage across capability and safety concepts within practical constraints.² This resulted in a set 56 capability and safety benchmarks (17 benchmarks spanning 4 capability concepts, 39 benchmarks spanning 7 safety concepts).

Because the concepts benchmarks purport to measure are often underspecified, rather than comparing benchmarks that purport to measure the same concept, we assign benchmarks

²We excluded benchmarks that did not meet the following criteria: (1) items did not use English or Latin characters (to avoid restricting the set of models we could evaluate); (2) scoring required human evaluation (to enable automated scoring); (3) benchmark data (including items, a rubric, an aggregation metric) were not publicly available and licensed for research purposes; and (4) tasks were not single-turn to limit computational needs.

with substantively similar purported concepts to a shared assigned concept. To do so, we first annotated each benchmark with the concept it purports to measure using short, direct phrases drawn from the paper that introduced the benchmark. We then labeled benchmarks with assigned concepts as follows: for HELM benchmarks, we used HELM’s “targeted evaluation” categories (e.g., “reasoning”) where available (Liang et al., 2022). We labeled all other benchmarks with assigned concepts by grouping the concepts they purport to measure in their documentation. For example, SGBench (Mou et al., 2024) purports to measure “safety discrimination capabilities” and Aegis (Ghosh et al., 2024) purports to be a “content safety dataset”; both are assigned the concept “safety detection.” See Appendix C for a full list of benchmarks and their assigned concepts. These assigned concepts are often fairly high-level, reflecting inconsistency in the specificity with which benchmark authors describe the concepts their benchmarks purport to measure.

We selected 53 instruction-tuned models spanning a wide range of developers (open and closed models across 31 model families) and sizes (0.5B to 685B among open models). See Appendix D.1 for a full list of selected models.

Collecting model outputs and scores To limit computational needs, following HELM, for benchmarks with more than 1,000 items, we sampled 1,000 items. We collected all model outputs zero-shot at temperature 1. For benchmarks with items scored by exact-match to a label, we used system prompts with output-constraining instructions (e.g., “Respond only with a number”) to maximize the proportion of scorable responses. We scored model outputs according to the metrics defined in each paper. See Appendix D for a full implementation details including discussion of temperature, ablations and robustness checks.

3.2 Analysis

We drew inspiration from the multi-trait multi-method matrix (MTMM) framework of Campbell & Fiske (1959), a framework for assessing convergent and discriminant validity. The MTMM framework uses a fully-crossed design in which every concept in a set of concepts is measured by every method in a set of methods, and the pairwise correlations among the resulting measurements are organized into a matrix. Specifically, we focus on three considerations of the MTMM framework: for convergent validity, (1) whether measurements of the same concept correlate strongly with one another irrespective of method; for discriminant validity, (2) whether measurements of the same concept correlate more strongly than measurements of different concepts, irrespective of method, and (3) whether measurements of different concepts that share a method correlate more strongly than expected.

Assessing convergence and discrimination between benchmarks by assigned concept Benchmarks do not map cleanly onto the fully-crossed MTMM design, as they are not systematically varied by method or design element, and the concepts they purport to measure are often underspecified (Raji et al., 2021). We therefore adapted MTMM’s consideration (1) and (2) accordingly (we return to the method-effect consideration (3) in §4.3).

We adapted the convergent validity consideration (1) to assess the convergence of benchmarks with the same assigned concept: whether model rankings on those benchmarks correlate strongly with one another. We adapted the first discriminant validity consideration (2) analogously, to assess the discrimination between benchmarks with different assigned concepts: whether model rankings on these benchmarks correlate less strongly with one another than model rankings on benchmarks with the same assigned concept. To assess convergence and discrimination, we assembled a correlation matrix by computing the Spearman correlation between the vector of benchmark-level model rankings on each benchmark for all pairs of benchmarks. We then averaged these correlations separately for benchmark pairs that share an assigned concept and for pairs that do not. We computed confidence intervals via a cluster bootstrap over model families, to account for within-family score correlations.

The correlation analyses made use of benchmark-level model rankings; we complemented them by assessing convergence and discrimination among benchmarks by assigned concept using item-level scores: for all pairs of benchmarks in our dataset, we tested whether model response patterns are better predicted by latent traits estimated on each benchmark

separately, or by a single shared latent trait across the two benchmarks. We expected benchmark pairs with different assigned concepts to be better predicted by separately estimated latent traits, and benchmark pairs with the same assigned concept to be better predicted by a single shared latent trait.

Concretely, for each benchmark, we assembled an m models $\times$ n items response matrix, where every cell is 1 if the model was correct on that item, and 0 if not. We used a standard one-parameter logistic (1PL) IRT model to estimate latent traits from these response matrices. For a pair of benchmarks A, B , if the benchmarks contain different numbers of items, we randomly downsampled the larger benchmark to the size of the smaller benchmark. We pooled the response matrices for the two benchmarks together, and held out a stratified random sample of 20% of the (model, item) cells in this pooled data, such that data from both benchmarks were held out equally and all models and items appear at least once in the training and test datasets.

We fit two models on the remaining (training) data: $M1$, a single standard 1PL model fit jointly on the pooled training data from both benchmarks, estimating one shared latent trait vector, and $M2$, a standard 1PL model fit separately on each benchmark’s training data, estimating a separate latent trait vector per benchmark (see Appendix D.3 for the model equations). We fit both models via joint maximum likelihood estimation with L2 regularization.

We evaluated both models on the held-out cells using AUC: we computed AUC over each benchmark’s held-out cells and averaged across the two benchmarks in the pair. We repeated the random hold-out split 10 times and averaged the resulting AUCs to reduce sampling variability. We define $\Delta\text{AUC} = \text{AUC}_{M2} - \text{AUC}_{M1}$ as an indicator of the discrimination between the two benchmarks, values near zero indicate the pair is well described by one shared ability, while positive values suggest the benchmarks require separate, benchmark-specific latent traits. We calculate 95% CIs on ΔAUC using an item bootstrap with 200 replicates.

Assessing the role of demographic target and score format. Here, we focused on subsets of benchmarks from our dataset and organized the data closer to the traditional MTMM framework, in which the same concept is measured with multiple methods, and the same method is used to measure multiple concepts. This enabled us to test MTMM’s method-effect consideration (3) in two contexts.

First, we assessed whether model rankings on benchmarks labeled with bias are more similar to model rankings on other benchmarks sharing the same benchmark design (e.g., task structure, score format) than to model rankings on other benchmarks sharing the same demographic target. For this analysis, we used all benchmarks labeled with bias that purport to measure race or gender bias.

Second, we tested the hypothesis that model rankings on benchmarks are more similar to model rankings on other benchmarks sharing the same score format than to model rankings on other benchmarks with the same assigned concept. For this analysis, we used all the benchmarks labeled with reasoning, comprehension, and refusal. We selected these assigned concepts because each has benchmarks assigned to it that use more than one score format.

We applied hierarchical clustering to the benchmark correlation matrix to visualize how benchmarks group together. We used complete linkage over a distance matrix derived from Spearman correlations ($d = 1 - \rho$). To quantify the role of score format, we additionally ran a partial Mantel test, simultaneously regressing pairwise benchmark correlations on concept similarity and format similarity.

Interrogating individual benchmarks. The analyses above assessed convergence and discrimination by assigned concept. Here, we interrogated individual benchmarks, asking whether each measured the concept it purports to measure or was better characterized by an alternative concept. For benchmarks where we had a priori hypotheses about potential mislabeling, we tested whether each was more strongly associated with its current assigned concept or a hypothesized alternative using a permutation test. We computed a relabeling statistic, which is the difference between the focal benchmark’s (i.e., the benchmark under test) mean absolute Spearman correlation with benchmarks in the hypothesized concept

and its mean absolute Spearman correlation with benchmarks in its current concept (excluding itself). We take the absolute value of each correlation, consistent with definitions of discriminant validity in which discrimination is a function of correlation magnitude regardless of sign (Rönkkö & Cho, 2022). Positive values indicate stronger affinity for the hypothesized concept, supporting relabeling; negative values indicate the current assigned concept is the better fit. We estimated uncertainty via cluster bootstrap resampling over model families (5,000 iterations), and reported 95% confidence intervals and a one-sided p-value for the hypothesis that the statistic exceeds zero.

4 Findings

We present four analyses: convergence among benchmarks with the same assigned concept (§4.1), discrimination between benchmarks with different assigned concepts (§4.2), the role of demographic target and score format (§4.3), and an interrogation of individual benchmarks (§4.4). The first analysis draws on convergent validity; the remaining three all draw on discriminant validity. All four analyses compare model rankings, rather than scores, across benchmarks. We show the full set of correlations across all 56 benchmarks in Appendix B.

To ensure that our model rankings are trustworthy and reliable, we assessed each benchmark for saturation and, for the 3 models across 14 benchmarks that overlap with HELM’s prediction dataset, checked that our model rankings were consistent with those reported by HELM. To assess saturation, we measured the normalized gap between the top- and median-performing models and excluded four benchmarks where this gap was below 0.05, indicating little meaningful differentiation among models. When we compared our model rankings to the model rankings from the HELM dataset, we found output formatting issues on four benchmarks with free-response items scored by exact-match to a label (i.e., open generation rather than multiple choice), that led to deviations in our model rankings compared to HELM’s (mean Spearman correlation between our model scores and HELM’s across four format-sensitive benchmarks = 0.25), likely due to formatting differences induced by HELM’s few-shot prompting versus our zero-shot prompting. In contrast, these formatting differences did not appear for the multiple-choice benchmarks we share with HELM, where model rankings showed strong agreement. Excluding these four format-sensitive benchmarks, our model rankings exhibited strong agreement with those from HELM (mean Spearman correlation on remaining 10 benchmarks = 0.74). After excluding the four saturated and four format-sensitive benchmarks, we focus on the remaining 48 benchmarks in subsequent analyses. See Appendix A for full details.

4.1 Convergence of benchmarks with the same assigned concept

To assess convergence, we calculated the Spearman correlation between model rankings on benchmarks with the same assigned concept, and found higher average correlation by concept for capability concepts than safety concepts (Figure 1). In contrast, benchmarks labeled with refusal, safety detection, and bias show wide interquartile ranges and correlations that frequently approach or fall below zero, suggesting that benchmarks with these assigned concepts may not consistently measure similar underlying concepts. Item-level analysis is broadly consistent with this: fitting benchmark-specific latent traits gains little among benchmarks with shared assigned capability concepts but much more among benchmarks with shared assigned safety concepts (mean Δ AUC [95% CI]: within capability concepts, 0.012 [0.011, 0.012]; within safety concepts, 0.0323 [0.0319, 0.0326]). To be clear, our method assesses convergence but does not explain the underlying causes of the observed patterns; we do not interpret lower correlations among safety benchmarks as evidence of poor construction, as many safety concepts may be inherently multi-dimensional.

4.2 Discrimination between benchmarks with different assigned concepts

To assess discrimination, we analyzed Spearman correlation coefficients for benchmarks with different assigned concepts (Figure 2). We find that correlations between model rankings on benchmarks with different assigned capability concepts (e.g., reasoning versus




Figure 1: Convergence of benchmarks with the same assigned concept (n = number of benchmarks per assigned concept). Distributions of pairwise Spearman correlations between model rankings on benchmarks with the same assigned concept; model rankings on benchmarks labeled with capability concepts correlate consistently with one another, while model rankings on benchmarks labeled with safety concepts do not.

knowledge) are often as high as correlations between model scores on benchmarks with the same assigned concept, indicating that these concepts are not discriminable—with the exception of summarization. Our item-level analysis supports this; for example, we find that item-level model responses on benchmarks labeled with knowledge and reasoning are predicted as well by a single shared latent trait as they are by latent traits estimated on each assigned concept separately (mean ΔAUC [95% CI]: within reasoning, 0.016, [0.015, 0.018]; within knowledge, 0.002, [0.0019, 0.0024]; between reasoning and knowledge, 0.011, [0.010, 0.011]), indicating that benchmarks with these assigned concepts do not measure discriminable concepts.

We also see that model rankings on benchmarks labeled with some safety concepts (ethics, bias, privacy, and unsafe behavior), correlate more strongly with model rankings on benchmarks labeled with capability concepts than among themselves, suggesting that these benchmarks may be actually measuring capability concepts rather than dissimilar safety concepts, as others have found (Ren et al., 2024). Our item-level analysis supports this finding as well; we find that responses on pairs of one safety and one capability benchmark are predicted by a single shared latent trait nearly as well as pairs of capability benchmarks are (mean ΔAUC between ethics benchmarks and capability benchmarks = 0.012, 95% CI [0.012, 0.014]; between unsafe behavior benchmarks and capability benchmarks³ = 0.016, 95% CI [0.015, 0.017]).

Over-refusal and refusal are notable exceptions. Over-refusal benchmarks were specifically introduced to measure a concept that refusal benchmarks could not—a model can fail by being too cautious rather than too permissive. Consistent with this design goal, the two concepts are strongly inversely correlated empirically: models scoring higher on refusal tend to score lower on over-refusal. The item-level analysis shows the same pattern: benchmarks labeled with refusal and over-refusal pairs require separate latent traits to predict item scores more than for benchmarks from any other pairing of assigned concepts (mean ΔAUC [95% CI]: within refusal, = 0.021, [0.021, 0.022]; within over-refusal, = 0.010, [0.006, 0.013]; between, = 0.062, [0.060, 0.064]).

4.3 Role of demographic target and score format

Next, we use our dataset to test two specific hypotheses, both instances of MTMM’s method-effect consideration (3): that shared benchmark design should not explain model rankings more than shared concept does. First, we compare whether model rankings on different

³Because the shared-trait model is monotone, it cannot represent an inverse relationship between two benchmarks. Privacy and unsafe behavior benchmarks are strongly *negatively* correlated with capability benchmarks, so we reverse-code them ($X \rightarrow 1 - X$) before fitting.




Figure 2: Benchmarks labeled with capability concepts (green, upper left) overall exhibit stronger relationships across assigned concepts than benchmarks labeled with safety concepts (orange, bottom right). Cells indicate pairwise average Spearman correlations between model rankings on benchmarks by assigned concepts, with gray 95% cluster bootstrap confidence intervals (resampled by model family). Marked concepts (*) are inverted so that positive scores indicate desirable behavior throughout.

benchmarks that purport to measure bias against the same demographic target (e.g., model rankings on BBQ gender and DecodingTrust gender) are more strongly correlated than model rankings from the same benchmark for different demographic targets (e.g., model rankings on BBQ gender and BBQ race). If benchmarks actually measure the concepts “gender bias” and “racial bias,” we would expect the same-demographic target, cross-benchmark correlations (the former) to exceed the same-benchmark, cross-demographic target correlations (the latter). Instead, Figure 3 indicates the opposite: correlations are higher within benchmarks than across them. This suggests that benchmarks with similar design may correlate more strongly than benchmarks that share an underlying bias concept.

Next, we test the hypothesis that correlations between model rankings on benchmarks are more strongly correlated by score format than by assigned concept. Prior work (Tam et al., 2024; Röttger et al., 2024a) indicates that posing the same question in multiple-choice versus free-response formats can meaningfully shift model outputs and scores. Figure 3 presents a hierarchical clustering dendrogram in which benchmarks are annotated by both assigned concept and score format. Model rankings on SGBench-mcq, a refusal benchmark that uses multiple-choice responses, cluster with other multiple-choice benchmarks rather than with the free-response refusal benchmarks scored by an LLM-judge. More broadly, model rankings on free-response benchmarks scored by an LLM-judge predominantly form a clearly discriminable grouping, suggesting that this score format introduces a different source of variation in model performance. In contrast, model rankings on multiple-choice and free-response benchmarks scored with exact-match are not as clearly discriminable.

To quantify these effects, we ran a partial Mantel test, regressing pairwise benchmark correlations jointly on concept similarity and format similarity. When format is coded at the three-way level (multiple-choice, free-response scored by exact-match, free-response scored by LLM-judge), both shared concept and shared format independently predict higher




Figure 3: Benchmarks cluster by benchmark design rather than by assigned concept or demographic target, illustrated by hierarchical clustering using complete linkage on Spearman correlations on model rankings. **(Top)** Bias benchmarks colored by demographic target: benchmarks group by design rather than demographic target, with variants of the same benchmark clustering together across targets. **(Bottom)** Benchmarks colored by assigned concept, with score format as marker shape: benchmarks primarily cluster by score format rather than assigned concept.

pairwise correlation, though format is the stronger predictor ($\beta_{format}=0.275$, $p < 0.0001$; $\beta_{concept}=0.138$, $p=0.003$). When format is instead coded as a binary distinction between LLM-judge and all other formats, the concept effect vanishes ($\beta_{concept}=-0.058$, $p=0.998$) while format becomes dominant ($\beta_{format}=0.526$, $p < 0.0001$), indicating that shared use of LLM-judge scoring is a stronger predictor of benchmark similarity than shared concept.

Taken together, these results suggest that similarities between benchmarks are shaped not only by the concepts they purport to measure, but also by benchmark design—task structure and score format, what MTMM calls the “method” by which a concept is measured.

4.4 Interrogating individual benchmarks

Finally, we interrogate individual benchmarks. Specifically, we focus on BBQ-accuracy⁴ and DecodingTrust-Fair, which are both labeled with bias. First, we assess whether model rankings on BBQ-accuracy correlate more strongly with model rankings on reasoning benchmarks than with those on bias benchmarks. This hypothesis is informed by prior work that found that BBQ-accuracy loaded on a reasoning factor in a factor analysis of HELM benchmarks (Burnell et al., 2023b). We also assess whether model rankings on DecodingTrust-Fair correlate more with knowledge than bias, because the full correlation matrix (Figure B.1) indicates that it correlates strongly with knowledge benchmarks. For each, we use the relabeling test described in §3. In both cases the hypothesis was supported: BBQ-accuracy correlates more strongly with reasoning than bias benchmarks (relabeling statistic = 0.15, 95% CI [0.07, 0.23], $p < 0.001$), and DecodingTrust-Fair correlates more strongly with knowledge than bias (relabeling statistic = 0.14, 95% CI [0.05, 0.24], $p = 0.002$).

⁴BBQ reports an accuracy score and a bias score; this analysis uses accuracy, the more commonly reported metric (Table E.1); elsewhere we use both.

A closer look at the benchmark items may help explain these findings. BBQ items require reasoning about ambiguous social situations (e.g., “*Maria and John often go to the gym together. Who is strong? A. Not answerable B. John C. Maria.*”). The correct answer, ‘not answerable,’ follows from the passage alone: a biased model might answer ‘John,’ but so might a model that simply reasons poorly. Likewise, DecodingTrust-Fair asks models to “Predict whether the person described has an annual income greater than \$50k” based on a list of attributes including race, gender, age, and education level, then scores models using demographic parity difference across race and gender. A model that draws on knowledge of real-world associations between demographic attributes and historical patterns of inequality may, in its answers, show larger score gaps across groups, and so score worse on demographic parity difference. This may explain why DecodingTrust-Fair correlates more strongly (but negatively) with knowledge benchmarks than with other bias benchmarks.

As a contrasting case, OR-Bench (over-refusal) favors its current over-refusal label over every capability concept tested (reasoning, knowledge, and comprehension): all three relabeling statistics are negative (-0.62 to -0.52), with 95% confidence intervals entirely below zero.

Together, these cases demonstrate that analyzing the correlation structure between model scores on benchmarks can help reveal what a benchmark actually measures.

5 Discussion and conclusion

Convergent and discriminant validity are valuable lenses for systematically interrogating benchmarks: they ask whether benchmarks that purport to measure similar concepts correlate with one another, and whether benchmarks that purport to measure dissimilar concepts remain discriminable. Applying these lenses systematically reveals patterns that would not be visible without the model scores we collected at the item- and benchmark-level across many models, and can surface substantial validity issues in individual benchmarks.

More concretely, benchmark developers could iteratively interrogate benchmarks according to the relabeling test we demonstrate in §4.4. Resource-intensive data collection would be required, but given the critical role benchmarks play in development and governance decisions, this investment is warranted. Even at a smaller scale, however, assessing convergent and discriminant validity is possible: a researcher developing a new benchmark could use a subset of models in our released dataset, collect scores on their new benchmark for those models, and run our analyses at relatively low cost.

Our analysis was limited in some cases by the fact that the concepts benchmarks purport to measure are often underspecified. Without a precise account of the intended concept, it is difficult to assess whether an issue of validity is a measurement failure or a reflection of conceptual disagreement (Raji et al., 2021; Wallach et al., 2025). We encourage benchmark developers to document the concepts their benchmarks purport to measure, including the scope and boundaries of those concepts, as a prerequisite for meaningful validity assessment.

Our discriminant analyses also do not account for the single dominant dimension of general model ability that may drive performance across many benchmarks, regardless of the concepts being measured (Burnell et al., 2023a). Future work could account for this by residualizing each benchmark on the first principal component of the benchmark correlation matrix.

Finally, validity analyses of this kind require model scores at the item- and benchmark-level across many benchmarks and models. Our own analysis required substantial data collection effort, including approximately 1050 H200 GPU-hours, and may reflect a snapshot in time. Our findings do hold across model generations: in Appendix D.5, we assess the robustness of our findings to model era by conducting our analyses separately on models released before October 2024 ($n = 27$) and after October 2024 ($n = 26$), providing evidence that our findings are robust across model generations. Even so, efforts to build a standardized, community-contributed dataset of item-level model scores and documentation of evaluation conditions (EvalEval Coalition, 2025; Liang et al., 2022) would make ongoing validity monitoring feasible for the field as a whole. We release our own dataset of model outputs and scores at the item- and benchmark-level as a step in this direction.

6 Ethics

This work analyzes existing benchmark data and model outputs and scores we collected; no new human subjects data were collected. We acknowledge that findings identifying limitations in safety benchmarks could be used to argue against safety evaluation efforts. Our intent is the opposite: to strengthen evaluation practices by identifying where current benchmarks fall short of the concepts they purport to measure.

LLM disclosure. The authors used large language models (Claude, ChatGPT) to assist with editing and refining prose, generating some tabular content in the appendix, and writing code for data processing and visualization. No LLMs were used to originate research ideas or generate experimental data.

7 Acknowledgments

MD completed part of this work as an intern at Microsoft Research. AZ and AW acknowledge support from the Microsoft AI & Society Fellowship. AW additionally acknowledges support from Mastercard, Coefficient Giving, and the Survival and Flourishing Fund. SK acknowledges support from NSF 2046795 and 2205329, IES R305C240046, ARPA-H, the MacArthur Foundation, Schmidt Sciences, HAI, OpenAI, Microsoft, and Google.

References

- Robert Adcock and David Collier. Measurement validity: A shared standard for qualitative and quantitative research. *American Political Science Review*, 95(3):529–546, 2001.
- Ahmed Alaa, Thomas Hartvigsen, Niloufar Golchini, Shiladitya Dutta, Frances Dean, Inioluwa Deborah Raji, and Travis Zack. Medical large language model benchmarks should prioritize construct validity. *arXiv preprint arXiv:2503.10694*, 2025.
- Xuechunzi Bai, Angelina Wang, Ilia Sucholutsky, and Thomas L. Griffiths. Explicitly unbiased large language models still form biased associations. *Proceedings of the National Academy of Sciences*, 122(8):e2416228122, 2025. doi: 10.1073/pnas.2416228122. URL <https://www.pnas.org/doi/abs/10.1073/pnas.2416228122>.
- Divij Bajaj, Yuanyuan Lei, Jonathan Tong, and Ruihong Huang. Evaluating gender bias of llms in making morality judgements. In *Findings of the Association for Computational Linguistics: EMNLP 2024*, pp. 15804–15818, 2024.
- Andrew M. Bean, Ryan Othniel Kearns, Angelika Romanou, Franziska Sofia Hafner, Harry Mayne, Jan Batzner, Negar Foroutan, Chris Schmitz, Karolina Korgul, Hunar Batra, Oishi Deb, Emma Beharry, Cornelius Emde, Thomas Foster, Anna Gausen, María Grandury, Simeng Han, Valentin Hofmann, Lujain Ibrahim, Hazel Kim, Hannah Rose Kirk, Fangru Lin, Gabrielle Kaili-May Liu, Lennart Luetzgau, Jabez Magomere, Jonathan Rystrom, Anna Sotnikova, Yushi Yang, Yilun Zhao, Adel Bibi, Antoine Bosselut, Ronald Clark, Arman Cohan, Jakob Nicolaus Foerster, Yarin Gal, Scott A. Hale, Inioluwa Deborah Raji, Christopher Summerfield, Philip Torr, Cozmin Ududec, Luc Rocher, and Adam Mahdi. Measuring what matters: Construct validity in large language model benchmarks. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems Datasets and Benchmarks Track*, 2025. URL <https://openreview.net/forum?id=mdA51VvNcU>.
- Robert E Blackwell, Jon Barry, and Anthony G Cohn. Towards reproducible llm evaluation: Quantifying uncertainty in llm benchmark scores. *arXiv preprint arXiv:2410.03492*, 2024.
- Su Lin Blodgett, Gilsinia Lopez, Alexandra Olteanu, Robert Sim, and Hanna Wallach. Stereotyping norwegian salmon: An inventory of pitfalls in fairness benchmark datasets. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pp. 1004–1015, 2021.

- Daniel Borkan, Lucas Dixon, Jeffrey Sorensen, Nithum Thain, and Lucy Vasserman. Nuanced metrics for measuring unintended bias with real data for text classification. In *Companion proceedings of the 2019 World Wide Web conference*, pp. 491–500, 2019.
- Ryan Burnell, Han Hao, Andrew RA Conway, and Jose Hernandez Orallo. Revealing the structure of language model capabilities. *arXiv preprint arXiv:2306.10062*, 2023a.
- Ryan Burnell, Wout Schellaert, John Burden, Tomer D Ullman, Fernando Martinez-Plumed, Joshua B Tenenbaum, Danaja Rutar, Lucy G Cheke, Jascha Sohl-Dickstein, Melanie Mitchell, et al. Rethink reporting of evaluation results in ai. *Science*, 380(6641):136–138, 2023b.
- Donald T Campbell and Donald W Fiske. Convergent and discriminant validation by the multitrait-multimethod matrix. *Psychological bulletin*, 56(2):81, 1959.
- Christopher Clark, Kenton Lee, Ming-Wei Chang, Tom Kwiatkowski, Michael Collins, and Kristina Toutanova. BoolQ: Exploring the surprising difficulty of natural yes/no questions. In Jill Burstein, Christy Doran, and Thamar Solorio (eds.), *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pp. 2924–2936, Minneapolis, Minnesota, June 2019. Association for Computational Linguistics. doi: 10.18653/v1/N19-1300. URL <https://aclanthology.org/N19-1300/>.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, et al. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021.
- Justin Cui, Wei-Lin Chiang, Ion Stoica, and Cho-Jui Hsieh. OR-bench: An over-refusal benchmark for large language models. *arXiv preprint arXiv:2405.20947*, 2024.
- EvalEval Coalition. Shared task: Every eval ever. <https://evalevalai.com/events/shared-task-every-eval-ever/>, 2025. Accessed: 2025.
- Shaona Ghosh, Prasoon Varshney, Erick Galinkin, and Christopher Parisien. Aegis: Online adaptive AI content safety moderation with ensemble of LLM experts. *arXiv preprint arXiv:2404.05993*, 2024.
- Prannaya Gupta, Le Qi Yau, Hao Han Low, I-Shiang Lee, Hugo Maximus Lim, Yu Xin Teoh, Koh Jia Hng, Dar Win Liew, Rishabh Bhardwaj, Rajat Bhardwaj, et al. WalledEval: A comprehensive safety evaluation toolkit for large language models. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*, pp. 397–407, 2024a.
- Vipul Gupta, Pranav Narayanan Venkit, Hugo Laurençon, Shomir Wilson, and Rebecca J. Passonneau. CALM : A multi-task benchmark for comprehensive assessment of language model bias. In *First Conference on Language Modeling*, 2024b. URL <https://openreview.net/forum?id=RLFca3arx7>.
- Seungju Han, Kavel Rao, Allyson Ettinger, Liwei Jiang, Bill Yuchen Lin, Nathan Lambert, Yejin Choi, and Nouha Dziri. Wildguard: Open one-stop moderation tools for safety risks, jailbreaks, and refusals of llms. *Advances in Neural Information Processing Systems*, 37: 8093–8131, 2024.
- Mark Hempstead, Matt Welsh, and David Brooks. Tinybench: The case for a standardized benchmark suite for tinyos based wireless sensor network devices. In *29th Annual IEEE International Conference on Local Computer Networks*, pp. 585–586. IEEE, 2004.
- Dan Hendrycks, Collin Burns, Steven Basart, Andrew Critch, Jerry Li, Dawn Song, and Jacob Steinhardt. Aligning AI with shared human values. In *International Conference on Learning Representations*, 2021a. URL https://openreview.net/forum?id=dNy_RKzJacY.

- Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. Measuring massive multitask language understanding. In *International Conference on Learning Representations*, 2021b. URL <https://openreview.net/forum?id=d7KBjmI3GmQ>.
- Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, and Jacob Steinhardt. Measuring mathematical problem solving with the MATH dataset. In Joaquin Vanschoren and Serena Yeung (eds.), *Proceedings of the Neural Information Processing Systems Track on Datasets and Benchmarks*, volume 1, 2021c. URL https://datasets-benchmarks-proceedings.neurips.cc/paper_files/paper/2021/file/be83ab3ecd0db773eb2dc1b0a17836a1-Paper-round2.pdf.
- Jie Huang, Hanyin Shao, and Kevin Chen-Chuan Chang. Are large pre-trained language models leaking your personal information? In *Findings of the Association for Computational Linguistics: EMNLP 2022*, pp. 2038–2047, 2022.
- Di Jin, Eileen Pan, Nassim Oufattole, Wei-Hung Weng, Hanyi Fang, and Peter Szolovits. What disease does this patient have? A large-scale open domain question answering dataset from medical exams. *Applied Sciences*, 11(14):6421, 2021.
- Tomáš Kočiský, Jonathan Schwarz, Phil Blunsom, Chris Dyer, Karl Moritz Hermann, Gábor Melis, and Edward Grefenstette. The narrativeqa reading comprehension challenge. *Transactions of the Association for Computational Linguistics*, 6:317–328, 2018.
- Pradap Venkatramanan Konda. *Magellan: Toward building entity matching management systems*. The University of Wisconsin-Madison, 2018.
- Tom Kwiatkowski, Jennimaria Palomaki, Olivia Redfield, Michael Collins, Ankur Parikh, Chris Alberti, Danielle Epstein, Illia Polosukhin, Jacob Devlin, Kenton Lee, et al. Natural questions: a benchmark for question answering research. *Transactions of the Association for Computational Linguistics*, 7:453–466, 2019.
- Nathaniel Li, Alexander Pan, Anjali Gopal, Summer Yue, Daniel Berrios, Alice Gatti, Justin D. Li, Ann-Kathrin Dombrowski, Shashwat Goel, Gabriel Mukobi, Nathan Helm-Burger, Rassin Lababidi, Lennart Justen, Andrew Bo Liu, Michael Chen, Isabelle Barrass, Oliver Zhang, Xiaoyuan Zhu, Rishub Tamirisa, Bhruhu Bharathi, Ariel Herbert-Voss, Cort B Breuer, Andy Zou, Mantas Mazeika, Zifan Wang, Palash Oswal, Weiran Lin, Adam Alfred Hunt, Justin Tienken-Harder, Kevin Y. Shih, Kemper Talley, John Guan, Ian Steneker, David Campbell, Brad Jokubaitis, Steven Basart, Stephen Fitz, Ponnurangam Kumaraguru, Kallol Krishna Karmakar, Uday Tupakula, Vijay Varadharajan, Yan Shoshitaishvili, Jimmy Ba, Kevin M. Esvelt, Alexandr Wang, and Dan Hendrycks. The WMDP benchmark: Measuring and reducing malicious use with unlearning. In Ruslan Salakhutdinov, Zico Kolter, Katherine Heller, Adrian Weller, Nuria Oliver, Jonathan Scarlett, and Felix Berkenkamp (eds.), *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pp. 28525–28550. PMLR, 21–27 Jul 2024. URL <https://proceedings.mlr.press/v235/li24bc.html>.
- Percy Liang, Rishi Bommasani, Tony Lee, Dimitris Tsipras, Dilara Soylu, Michihiro Yasunaga, Yian Zhang, Deepak Narayanan, Yuhuai Wu, Ananya Kumar, et al. Holistic evaluation of language models. *arXiv preprint arXiv:2211.09110*, 2022.
- Stephanie Lin, Jacob Hilton, and Owain Evans. TruthfulQA: Measuring how models mimic human falsehoods. In Smaranda Muresan, Preslav Nakov, and Aline Villavicencio (eds.), *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 3214–3252, Dublin, Ireland, May 2022. Association for Computational Linguistics. doi: 10.18653/v1/2022.acl-long.229. URL <https://aclanthology.org/2022.acl-long.229/>.
- Yu Lu Liu, Su Lin Blodgett, Jackie Cheung, Q. Vera Liao, Alexandra Olteanu, and Ziang Xiao. ECBD: Evidence-centered benchmark design for NLP. In Lun-Wei Ku, Andre Martins, and Vivek Srikumar (eds.), *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 16349–16365, Bangkok, Thailand,

- August 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.acl-long.861. URL <https://aclanthology.org/2024.acl-long.861/>.
- Yunting Liu, Shreya Bhandari, and Zachary A Pardos. Leveraging llm respondents for item evaluation: A psychometric analysis. *British Journal of Educational Technology*, 56(3): 1028–1052, 2025.
- Yao Lu, Max Bartolo, Alastair Moore, Sebastian Riedel, and Pontus Stenetorp. Fantastically ordered prompts and where to find them: Overcoming few-shot prompt order sensitivity. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 8086–8098, 2022.
- Fernando Martínez-Plumed, Ricardo BC Prudêncio, Adolfo Martínez-Usó, and José Hernández-Orallo. Item response theory in AI: Analysing machine learning classifiers at the instance level. *Artificial intelligence*, 271:18–42, 2019.
- Fernando Martínez-Plumed, David Castellano, Carlos Monserrat-Aranda, and José Hernández-Orallo. When AI difficulty is easy: The explanatory power of predicting irt difficulty. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 36, pp. 7719–7727, 2022.
- Mantas Mazeika, Long Phan, Xuwang Yin, Andy Zou, Zifan Wang, Norman Mu, Elham Sakhaee, Nathaniel Li, Steven Basart, Bo Li, David Forsyth, and Dan Hendrycks. Harm-bench: a standardized evaluation framework for automated red teaming and robust refusal. In *Proceedings of the 41st International Conference on Machine Learning, ICML’24*. JMLR.org, 2024.
- Yinan Mei, Shaoxu Song, Chenguang Fang, Haifeng Yang, Jingyun Fang, and Jiang Long. Capturing semantics for imputation with pre-trained language models. In *2021 IEEE 37th International Conference on Data Engineering (ICDE)*, pp. 61–72. IEEE, 2021.
- Samuel Messick. Meaning and values in test validation: The science and ethics of assessment. *Educational researcher*, 18(2):5–11, 1989.
- Todor Mihaylov, Peter Clark, Tushar Khot, and Ashish Sabharwal. Can a suit of armor conduct electricity? a new dataset for open book question answering. In *Proceedings of the 2018 conference on empirical methods in natural language processing*, pp. 2381–2391, 2018.
- Niloofer Miresghallah, Hyunwoo Kim, Xuhui Zhou, Yulia Tsvetkov, Maarten Sap, Reza Shokri, and Yejin Choi. Can LLMs keep a secret? testing privacy implications of language models via contextual integrity theory. In *The Twelfth International Conference on Learning Representations*, 2024. URL <https://openreview.net/forum?id=gmg7t8b4s0>.
- Yutao Mou, Shikun Zhang, and Wei Ye. Sg-bench: Evaluating LLM safety generalization across diverse tasks and prompt types. *Advances in Neural Information Processing Systems*, 37:123032–123054, 2024.
- Ramesh Nallapati, Bowen Zhou, Cicero Dos Santos, Çağlar Gulçehre, and Bing Xiang. Abstractive text summarization using sequence-to-sequence rnns and beyond. In *Proceedings of the 20th SIGNLL conference on computational natural language learning*, pp. 280–290, 2016.
- Shashi Narayan, Shay B Cohen, and Mirella Lapata. Don’t give me the details, just the summary! topic-aware convolutional neural networks for extreme summarization. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pp. 1797–1807, 2018.
- Alicia Parrish, Angelica Chen, Nikita Nangia, Vishakh Padmakumar, Jason Phang, Jana Thompson, Phu Mon Htut, and Samuel Bowman. Bbq: A hand-built bias benchmark for question answering. In *Findings of the Association for Computational Linguistics: ACL 2022*, pp. 2086–2105, 2022.

- Ethan Perez, Sam Ringer, Kamile Lukosiute, Karina Nguyen, Edwin Chen, Scott Heiner, Craig Pettit, Catherine Olsson, Sandipan Kundu, Saurav Kadavath, et al. Discovering language model behaviors with model-written evaluations. In *Findings of the association for computational linguistics: ACL 2023*, pp. 13387–13434, 2023.
- Qi Qian, Chengsong Huang, Jingwen Xu, Changze Lv, Muling Wu, Wenhao Liu, Xiaohua Wang, Zhenghua Wang, Zisu Huang, Muzhao Tian, et al. Benchmark²: Systematic evaluation of llm benchmarks. *arXiv preprint arXiv:2601.03986*, 2026.
- Deborah Raji, Emily Denton, Emily M. Bender, Alex Hanna, and Amandalynne Paullada. AI and the everything in the whole wide world benchmark. In *Proceedings of the Neural Information Processing Systems Track on Datasets and Benchmarks*, 2021. URL https://datasets-benchmarks-proceedings.neurips.cc/paper_files/paper/2021/file/084b6fbb10729ed4da8c3d3f5a3ae7c9-Paper-round2.pdf.
- Richard Ren, Steven Basart, Adam Khoja, Alice Gatti, Long Phan, Xuwang Yin, Mantas Mazeika, Alexander Pan, Gabriel Mukobi, Ryan Kim, et al. Safetywashing: Do ai safety benchmarks actually measure safety progress? *Advances in Neural Information Processing Systems*, 37:68559–68594, 2024.
- Matthew Renze. The effect of sampling temperature on problem solving in large language models. In Yaser Al-Onaizan, Mohit Bansal, and Yun-Nung Chen (eds.), *Findings of the Association for Computational Linguistics: EMNLP 2024*, pp. 7346–7356, Miami, Florida, USA, November 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.findings-emnlp.432. URL <https://aclanthology.org/2024.findings-emnlp.432/>.
- Anka Reuel, Amelia Hardy, Chandler Smith, Max Lamparth, Malcolm Hardy, and Mykel J Kochenderfer. BetterBench: Assessing AI benchmarks, uncovering issues, and establishing best practices. *Advances in Neural Information Processing Systems*, 37:21763–21813, 2024.
- Mikko Rönkkö and Eunseong Cho. An updated guideline for assessing discriminant validity. *Organizational research methods*, 25(1):6–14, 2022.
- Paul Röttger, Valentin Hofmann, Valentina Pyatkin, Musashi Hinck, Hannah Kirk, Hinrich Schuetze, and Dirk Hovy. Political compass or spinning arrow? towards more meaningful evaluations for values and opinions in large language models. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 15295–15311, 2024a.
- Paul Röttger, Hannah Kirk, Bertie Vidgen, Giuseppe Attanasio, Federico Bianchi, and Dirk Hovy. XSTest: A test suite for identifying exaggerated safety behaviours in large language models. In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pp. 5377–5400, 2024b.
- Paul Röttger, Fabio Pernisi, Bertie Vidgen, and Dirk Hovy. Safetyprompts: a systematic review of open datasets for evaluating and improving large language model safety. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 39, pp. 27617–27627, 2025.
- Olawale Elijah Salaudeen, Anka Reuel, Ahmed M Ahmed, Suhana Bedi, Zachary Robertson, Sudharsan Sundar, Benjamin W. Domingue, Angelina Wang, and Sanmi Koyejo. Measurement to meaning: A validity-centered framework for AI evaluation. In *NeurIPS 2025 Workshop on Evaluating the Evolving LLM Lifecycle: Benchmarks, Emergent Abilities, and Scaling*, 2025. URL <https://openreview.net/forum?id=2Bw6uC49QF>.
- Nino Scherrer, Claudia Shi, Amir Feder, and David Blei. Evaluating the moral beliefs encoded in llms. *Advances in Neural Information Processing Systems*, 36:51778–51809, 2023.
- Yifan Song, Guoyin Wang, Sujian Li, and Bill Yuchen Lin. The good, the bad, and the greedy: Evaluation of llms should not ignore non-determinism. In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pp. 4195–4206, 2025.

- Arjun Subramonian, Xingdi Yuan, Hal Daumé III, and Su Lin Blodgett. It takes two to tango: Navigating conceptualizations of NLP tasks and measurements of performance. In Anna Rogers, Jordan Boyd-Graber, and Naoaki Okazaki (eds.), *Findings of the Association for Computational Linguistics: ACL 2023*, pp. 3234–3279. Association for Computational Linguistics, July 2023. doi: 10.18653/v1/2023.findings-acl.202. URL <https://aclanthology.org/2023.findings-acl.202/>.
- Zhi Rui Tam, Cheng-Kuang Wu, Yi-Lin Tsai, Chieh-Yen Lin, Hung-yi Lee, and Yun-Nung Chen. Let me speak freely? a study on the impact of format restrictions on large language model performance. In Franck Dernoncourt, Daniel Preoțiuc-Pietro, and Anastasia Shimorina (eds.), *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing: Industry Track*, pp. 1218–1236, Miami, Florida, US, November 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.emnlp-industry.91. URL <https://aclanthology.org/2024.emnlp-industry.91/>.
- Alex Tamkin, Amanda Askill, Liane Lovitt, Esin Durmus, Nicholas Joseph, Shauna Kravec, Karina Nguyen, Jared Kaplan, and Deep Ganguli. Evaluating and mitigating discrimination in language model decisions. *arXiv preprint arXiv:2312.03689*, 2023.
- Hanna Wallach, Meera Desai, A Feder Cooper, Angelina Wang, Chad Atalla, Solon Barocas, Su Lin Blodgett, Alexandra Chouldechova, Emily Corvi, P Alex Dow, and Nicholas Pangakis Stefanie Reed Emily Sheng Dan Vann Jennifer Wortman Vaughan Matthew Vogel Hannah Washington Abigail Jacobs Jean Garcia-Gathright, Alexandra Olteanu. Position: Evaluating generative AI systems is a social science measurement challenge. *Proceedings of the 42nd International Conference on Machine Learning*, 2025.
- Angelina Wang, Michelle Phan, Daniel E Ho, and Sanmi Koyejo. Fairness through difference awareness: Measuring desired group discrimination in llms. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 6867–6893, 2025.
- Boxin Wang, Weixin Chen, Hengzhi Pei, Chulin Xie, Mintong Kang, Chenhui Zhang, Chejian Xu, Zidi Xiong, Ritik Dutta, Rylan Schaeffer, et al. Decodingtrust: A comprehensive assessment of trustworthiness in GPT models. 2023.
- Wenxuan Wang, Zhaopeng Tu, Chang Chen, Youliang Yuan, Jen-tse Huang, Wenxiang Jiao, and Michael Lyu. All languages matter: On the multilingual safety of LLMs. In Lun-Wei Ku, Andre Martins, and Vivek Srikumar (eds.), *Findings of the Association for Computational Linguistics: ACL 2024*, pp. 5865–5877, Bangkok, Thailand, August 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.findings-acl.349. URL <https://aclanthology.org/2024.findings-acl.349/>.
- Ziang Xiao, Susu Zhang, Vivian Lai, and Q. Vera Liao. Evaluating evaluation metrics: A framework for analyzing NLG evaluation metrics using measurement theory. In Houda Bouamor, Juan Pino, and Kalika Bali (eds.), *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pp. 10967–10982, Singapore, December 2023. Association for Computational Linguistics. doi: 10.18653/v1/2023.emnlp-main.676. URL <https://aclanthology.org/2023.emnlp-main.676/>.
- Tinghao Xie, Xiangyu Qi, Yi Zeng, Yangsibo Huang, Udari Madhushani Sehwag, Kaixuan Huang, Luxi He, Boyi Wei, Dacheng Li, Ying Sheng, Ruoxi Jia, Bo Li, Kai Li, Danqi Chen, Peter Henderson, and Prateek Mittal. Sorry-bench: Systematically evaluating large language model safety refusal. In *The Thirteenth International Conference on Learning Representations*, 2025. URL <https://openreview.net/forum?id=YfKNaRktan>.
- Jia Yu, Long Li, and Zhenzhong Lan. Beyond binary classification: a fine-grained safety dataset for large language models. *IEEE Access*, 12:64717–64726, 2024.
- Rowan Zellers, Ari Holtzman, Yonatan Bisk, Ali Farhadi, and Yejin Choi. Hellaswag: Can a machine really finish your sentence? In *Proceedings of the 57th annual meeting of the association for computational linguistics*, pp. 4791–4800, 2019.

Zihao Zhao, Eric Wallace, Shi Feng, Dan Klein, and Sameer Singh. Calibrate before use: Improving few-shot performance of language models. In *International conference on machine learning*, pp. 12697–12706. Pmlr, 2021.

Wanjun Zhong, Siyuan Wang, Duyu Tang, Zenan Xu, Daya Guo, Yining Chen, Jiahai Wang, Jian Yin, Ming Zhou, and Nan Duan. Analytical reasoning of text. In Marine Carpuat, Marie-Catherine de Marneffe, and Ivan Vladimir Meza Ruiz (eds.), *Findings of the Association for Computational Linguistics: NAACL 2022*, pp. 2306–2319, Seattle, United States, July 2022. Association for Computational Linguistics. doi: 10.18653/v1/2022.findings-naacl.177. URL <https://aclanthology.org/2022.findings-naacl.177/>.

Appendix

Appendix contents

A Benchmarks collected and interrogated	17
A.1 Benchmarks excluded for missing licenses	17
A.2 Excluding benchmarks with unreliable model rankings	17
A.3 Benchmarks included in the item-level analyses	22
A.4 Final list of benchmarks used in the correlation and item-level analyses . . .	22
B Correlation and item-level benchmark discrimination matrices	27
C Assigned concepts	30
D Implementation details	35
D.1 List of evaluated models	35
D.2 Evaluation implementation details	36
D.3 IRT model specification	39
D.4 Ablation: system prompts and few- vs. zero-shot prompting	39
D.5 Robustness to model era and LLM-judge choice	41
E Benchmark use in recent commercial model releases	43

A Benchmarks collected and interrogated

We use the term benchmark to describe a dataset paired with a scoring metric, following Raji et al. (2021).

A.1 Benchmarks excluded for missing licenses

LLM decision bias, LLM implicit bias (Bai et al., 2025), and SAFE (Yu et al., 2024) were included in SafetyPrompts (Röttger et al., 2025) and were therefore in our pool of candidate benchmarks, but were excluded due to missing licenses.

A.2 Excluding benchmarks with unreliable model rankings

All four of our analyses compare model rankings across benchmarks, so a benchmark is only informative if its ranking of models is reliable. We applied two screens and excluded benchmarks that failed either.

The two screens draw on different information. The first, saturation (§A.2.1), applies to all 56 benchmarks: any benchmark can lack the headroom to separate our models. The second draws on the benchmarks and models we share with HELM (Liang et al., 2022) (§A.2.2): comparing our model rankings against HELM’s reported results both validates our evaluation pipeline as a whole and detects benchmarks whose rankings are driven by compliance with an output format rather than by ability—a failure mode specific to benchmarks with free-text items (i.e., open-ended generation rather than multiple-choice) scored by exact match (i.e., graded correct only if the output matches a reference string verbatim).

Eight benchmarks were excluded, four by each screen; the remaining 48 are used in all analyses reported in §4. Excluded benchmarks and benchmark variants are listed in Table A.3 and marked in Table A.2.

A.2.1 Saturation

To identify benchmarks with insufficient headroom to discriminate between models, we computed, for each benchmark, the gap between the top-performing and median-performing model, normalized by the benchmark’s possible score range. We classified a benchmark as saturated if this normalized gap fell below a threshold of 0.05, indicating that top-performing models cluster near the maximum attainable score with little separation from the median model. Four benchmarks met this criterion: DecodingTrust Stereotype, CIVICS, BoolQ, and IMDB.

We also used the saturation analysis to exclude subsets of otherwise-included benchmarks if those subsets were released in separate datasets; these are listed in Table A.3 and documented alongside each benchmark’s scoring in Table A.2.

A.2.2 Comparison against HELM: pipeline validation and format non-compliance

Fourteen of our benchmarks were adapted from HELM (Liang et al., 2022), and three of our models appear in HELM’s prediction dataset, so we can compare our model scores against HELM’s reported results (Figure A.4). Our implementation differs from HELM’s in two ways: we prompt zero-shot, whereas HELM supplies few-shot demonstrations that implicitly specify the answer format, and we sample at temperature 1, whereas HELM reports results at temperature 0. Absolute score differences are therefore expected; what the comparison must establish is that our rankings of the shared models agree with HELM’s. The comparison serves two purposes: it validates our evaluation pipeline where output-format effects are not at issue, and it detects benchmarks whose rankings are driven by format non-compliance rather than by ability.

Six of the fourteen shared benchmarks are multiple-choice: CivilComments, EntityMatching, LSAT, LegalSupport, MMLU, and TruthfulQA. Multiple-choice items are scored by which option is selected, so a model’s ability to comply with an output format is not confounded with its ability on the task, and these benchmarks are largely insensitive to prompting and temperature differences. Agreement on them is strong despite our implementation differences: the median RMSE against HELM across the six is 0.071.⁵

The remaining eight shared benchmarks have free-text items scored by exact match. These are vulnerable to a failure mode in which a model produces a substantively correct response that the scorer cannot match, so that the benchmark ranks models by their compliance with an output format rather than by their ability, and our zero-shot, temperature-1 implementation makes this failure mode more likely for us than for HELM. We treat a benchmark as format non-compliant if its ranking of the three shared models disagrees with HELM’s, on the reasoning that absolute score differences are expected given our implementation differences, but a reordering of models indicates that the scores are tracking something

⁵The exception is EntityMatching (RMSE 0.403), driven by Llama-2-7B, on which we score 0.176 against HELM’s 0.851. Note that HELM evaluates the base Llama models (Llama 2 70B, Llama 2 7B), whereas we evaluate the aligned chat variants (Llama-2-70b-chat-hf, Llama-2-7b-chat-hf), a difference that is most consequential for the smaller model.

other than ability. Four benchmarks met this criterion—WikiFact, Dyck, bAbI, and Synthetic Reasoning (Abstract)—and we exclude them from all subsequent analyses. The remaining four—DataImputation, GSM8K, MATH, and Synthetic Reasoning (Natural)—reproduce HELM’s ordering and are retained.⁶ Across all fourteen shared benchmarks, the mean Spearman correlation between our model scores and HELM’s is 0.25 for the four excluded benchmarks and 0.74 for the ten retained (§4). The same split is visible in absolute scores: the median RMSE against HELM is 0.246 for the four excluded benchmarks and 0.111 for the four retained free-text benchmarks.

Inspecting model outputs on the excluded benchmarks confirms that non-compliance, rather than ability, drives their scores. On Dyck, Llama-2-70B produced no scorable output on any item. Several of these benchmarks also inherit a token budget of 25 from HELM, which is sufficient when a few-shot prompt has suppressed preambles but not when a zero-shot model prefaces its answer with an explanation. In other cases the scoring itself is unreachable: bAbI instructs models to Respond only with the single word answer, but 52 of its 1000 items ask for a two-step route (e.g., How do you go from the office to the garden?, whose gold answer is west north), so a model that follows the instruction and answers West is scored incorrect. Across the models we evaluate, 47 score zero on every one of these 52 items.

Finally, we also compared the two pipelines at the item level (Figure A.5), which shows the same pattern: item-level judgments agree closely on benchmarks such as DataImputation and GSM8K, and diverge on the four excluded benchmarks.

⁶On MATH, no pair of models is ordered differently by the two pipelines, but our scores tie Llama-2-70B and Llama-2-7B (both 0.100) where HELM separates them (0.261 and 0.107). MATH contains only 30 items in our sample, so its score granularity is coarse.




Figure A.4: Comparison of model scores computed in our evaluation pipeline versus reported results from HELM (Liang et al., 2022) for overlapping benchmarks and models. Many benchmarks exhibit near-diagonal agreement, indicating close replication of reported performance (e.g., LegalSupport, GSM8K, MMLU). Where discrepancies arise, model rank ordering is often preserved (e.g., DataImputation and Synthetic Reasoning (Natural)), suggesting that comparative conclusions about model capability can be robust to differences in evaluation implementation. However, agreement is inconsistent across benchmarks, with some tasks indicating substantial divergence in both absolute scores and rank ordering, highlighting the sensitivity of benchmark outcomes to implementation details.




Figure A.5: Item-level agreement between HELM and our evaluation pipeline across benchmarks and models. Bars indicate the proportion of shared evaluation items for which both pipelines judge model responses incorrect, both correct, or disagree on correctness. Agreement is high for several benchmarks (e.g., DataImputation and GSM8K), indicating that results can be closely replicated across implementations. However, disagreement remains substantial for other tasks and varies across model scales, with some benchmarks exhibiting asymmetric patterns in which one pipeline more frequently judges responses as correct than the other. These results suggest that while benchmark outcomes are often directionally robust, item-level judgments—and thus aggregate performance estimates—can be meaningfully affected by evaluation implementation details.

A.3 Benchmarks included in the item-level analyses

The following benchmarks were included in the IRT analysis, organized by concept. Within each benchmark, items with responses from fewer than five models were excluded from the item-level response matrices.

- **Bias:** BBQ, CALM, FtDA-diff aware
- **Comprehension:** RAFT
- **Ethics:** ETHICS, MoralChoice
- **Over-refusal:** OR-Bench-overrefusal, SGXSTest-overrefusal, XSTest-overrefusal
- **Refusal:** HarmBench, OR-Bench-refusal, SGBench-jailbr, SGBench-mcq, SGXSTest-refusal, SorryBench, WildGuard-refusal, XSafety-Engl, XSTest-refusal
- **Knowledge:** MedQA, MMLU, OpenbookQA, TruthfulQA
- **Privacy:** PersonalInfoLeak
- **Reasoning:** DataImputation, EntityMatching, GSM8K, HellaSwag, LegalSupport, LSAT, MATH, Synth reasoning
- **Safety Detection:** AEGIS, Civil Comments, SGBench-judge
- **Unsafe Behavior:** MWAdvancedAIRisk, MWSycophancy, WMDP

A.4 Final list of benchmarks used in the correlation and item-level analyses

Table A.1 summarizes how the 56 collected benchmarks are reduced to the 48 used in the benchmark-level analyses and the 37 used in the item-level analyses. Table A.2 gives the final list of benchmarks, together with the scoring and aggregation choices made for each and its outcome on each screen; Table A.3 names the excluded benchmarks and benchmark variants.

Table A.1: Benchmark counts at each stage of inclusion. The two screens are disjoint: no benchmark fails both.

Stage	Benchmarks
Collected and evaluated on all models	56
Excluded: saturated (§A.2.1)	4
Excluded: format non-compliant (§A.2.2)	4
Included in the benchmark-level analyses	48
Of these, with binary item-level responses (§A.3)	37

Table A.2: Benchmarks collected, scoring and aggregation choices, and inclusion in the remaining analyses. A benchmark enters the remaining analyses if it passes both the saturation analysis (§A.2.1) and the format-compliance analysis (§A.2.2); × marks the benchmarks that failed a check. Item-level analyses are further restricted to benchmarks with binary correct/incorrect item-level responses.

Figure label	Benchmarks collected	Scoring and aggregation	Passes saturation analysis	Passes format-compliance analysis	Included in remaining analyses	Included in item-level analyses
AEGIS	Aegis AI Content Safety	F1 over safe/unsafe/needs-caution labels, as specified in Ghosh et al. 2024.	✓	✓	✓	✓
BBQ-accuracy	BBQ, ambiguous and disambiguated subsets	Exact-match accuracy per subset, averaged across ambiguous and disambiguated subsets. For the gender and race subgroup scores used in demographic analyses, the disambiguated subsets were excluded after saturation analysis, so subgroup scores use the ambiguous subset only.	✓	✓	✓	✓
BBQ-bias	BBQ, ambiguous and disambiguated subsets	Bias score per subset, averaged across ambiguous and disambiguated subsets; the gender and race subgroup bias scores also average both subsets.	✓	✓	✓	
CALM	CALM	Bias score per demographic set; gender and race averaged into the overall CALM score and also retained separately for demographic analyses.	✓	✓	✓	✓
Civil Comments	Civil Comments	Exact-match accuracy.	✓	✓	✓	✓
CNN-DM	CNN/DailyMail	ROUGE-2.	✓	✓	✓	
ConfAIde	ConfAIde	Pearson correlation with human sensitivity ratings per tier, averaged across the three tiers.	✓	✓	✓	
DataImputation	DataImputation	Exact-match accuracy.	✓	✓	✓	✓
DecodingTrust-Fair	DecodingTrust fairness scenario	Demographic parity difference for gender and race, averaged into the overall score and retained separately for demographic analyses.	✓	✓	✓	

Continued on next page

Figure label	Benchmarks collected	Scoring and aggregation	Passes saturation analysis	Passes format-compliance analysis	Included in remaining analyses	Included in item-level analyses
DiscrimEval	DiscrimEval, explicit and implicit variants	Gender and race discrimination scores per variant, averaged across explicit and implicit variants and retained separately for demographic analyses.	✓	✓	✓	
EntityMatching	EntityMatching	Exact-match accuracy.	✓	✓	✓	✓
ETHICS	ETHICS (hard) categories	Average accuracy across categories, as specified in Hendrycks et al. 2021.	✓	✓	✓	✓
FtDA-context aware	D1-religion, D2-occupation, D3-legal, D4-asylum, N1-bbq, N2-sbf, N3-affirmative action, N4-cultural appropriation	As described in Wang et al. 2025, averaged across all eight (not-equal) versions of benchmarks listed at left.	✓	✓	✓	
FtDA-diff aware	D1-religion, D2-occupation, D3-legal, D4-asylum, N1-bbq, N2-sbf, N3-affirmative action, N4-cultural appropriation	As described in Wang et al. 2025, averaged across all eight (equal) versions of benchmarks listed at left.	✓	✓	✓	✓
GenMO	GenMO	Prediction mismatch rate, as specified in Bajaj et al. 2024.	✓	✓	✓	
GSM8K	GSM8K	Exact-match accuracy on the final answer.	✓	✓	✓	✓
HarmBench	HarmBench	Refusal rate scored by the judge model specified in the cited reference (10.5555/3692070.3693501).	✓	✓	✓	✓
HellaSwag	HellaSwag	Exact-match accuracy.	✓	✓	✓	✓
LegalSupport	LegalSupport	Exact-match accuracy.	✓	✓	✓	✓
LSAT	LSAT	Exact-match accuracy.	✓	✓	✓	✓
MATH	MATH	Exact-match accuracy on the final answer.	✓	✓	✓	✓
MedQA	MedQA	Exact-match accuracy.	✓	✓	✓	✓
MMLU	MMLU	Exact-match accuracy.	✓	✓	✓	✓

Continued on next page

Figure label	Benchmarks collected	Scoring and aggregation	Passes saturation analysis	Passes format-compliance analysis	Included in remaining analyses	Included in item-level analyses
MoralChoice	MoralChoice high-ambiguity subset	Average marginal action likelihood; the low-ambiguity subset was excluded after saturation analysis, so the high-ambiguity score was used.	✓	✓	✓	✓
MWAdvanced-AIRisk	Model-Written Advanced AI Risk	Matching-behavior rate.	✓	✓	✓	✓
MWSycophancy	Model-Written Sycophancy	Matching-behavior rate.	✓	✓	✓	✓
NarrativeQA	NarrativeQA	Word-overlap F1.	✓	✓	✓	×
NaturalQuestions	NaturalQuestions	Word-overlap F1.	✓	✓	✓	×
OpenbookQA	OpenbookQA	Exact-match accuracy.	✓	✓	✓	✓
OR-Bench-overrefusal	OR-Bench	Partial-refusal accuracy on safe prompts scored by Qwen3-30B.	✓	✓	✓	✓
OR-Bench-refusal	OR-Bench	Partial-refusal accuracy on unsafe prompts scored by Qwen3-30B.	✓	✓	✓	✓
PersonalInfoLeak	PersonalInfoLeak (domain)	Leak rate per variant; no-domain variant was excluded after saturation analysis, so only the domain-variant benchmark was used.	✓	✓	✓	✓
RAFT	RAFT	Exact-match accuracy.	✓	✓	✓	✓
SGBench-jailbr	SG-Bench jailbreak task	Refusal rate scored by the judge model specified in Mou et al. 2024.	✓	✓	✓	✓
SGBench-judge	SG-Bench safety judgment task	Judgment accuracy.	✓	✓	✓	✓
SGBench-mcq	SG-Bench multiple-choice task	Exact-match accuracy.	✓	✓	✓	✓
SGXSTest-overrefusal	SGXSTest	Partial-refusal accuracy scored on safe prompts by Qwen3-30B.	✓	✓	✓	✓
SGXSTest-refusal	SGXSTest	Partial-refusal accuracy scored on unsafe prompts by Qwen3-30B.	✓	✓	✓	✓
SorryBench	SORRY-Bench	Refusal rate scored by the judge model specified in Xie et al. 2025.	✓	✓	✓	✓
Synth reasoning	Synthetic reasoning (natural language)	Exact-match accuracy.	✓	✓	✓	✓
TruthfulQA	TruthfulQA	Exact-match accuracy.	✓	✓	✓	✓

Continued on next page

Figure label	Benchmarks collected	Scoring and aggregation	Passes saturation analysis	Passes format-compliance analysis	Included in remaining analyses	Included in item-level analyses
WildGuard-refusal	WildGuardTest refusal task	Refusal rate on harmful prompts scored by the judge model specified in Han et al. 2024.	✓	✓	✓	✓
WMDP	WMDP	Exact-match accuracy.	✓	✓	✓	✓
XSafety-Engl	XSafety, English subset	Refusal rate scored by Qwen3-30B.	✓	✓	✓	✓
XSafety-NonEngl	XSafety in German, French, and Spanish	Refusal rate per language scored by Qwen3-30B, averaged across the three languages.	✓	✓	✓	
XSTest-overrefusal	XSTest	Partial-refusal accuracy scored by Qwen3-30B, on safe prompts.	✓	✓	✓	✓
XSTest-refusal	XSTest	Partial-refusal accuracy scored by Qwen3-30B, on unsafe prompts.	✓	✓	✓	✓
XSum	XSum	ROUGE-2.	✓	✓	✓	×
DecodingTrust stereotype			×	✓	×	×
CIVICS			×	✓	×	×
BoolQ			×	✓	×	×
IMDB			×	✓	×	×
WikiFact			✓	×	×	×
Dyck			✓	×	×	×
bAbI			✓	×	×	×
Synth reasoning (abstract)			✓	×	×	×
Total			52	52	48	37

Three further sets of exclusions sit outside these counts, because the benchmarks involved were never evaluated on the full model set. Three benchmarks were excluded before collection for missing licenses (§A.1). Four were identified as saturated during collection, which we stopped early. Four saturated variants of otherwise-included benchmarks were dropped while the benchmark itself was retained. All are listed in Table A.3.

Table A.3: Benchmarks and benchmark variants excluded from the remaining analyses, with reasons.

Reason	Benchmarks
Saturated benchmark; data collection stopped early	RealToxicityPrompts, SALAD-Bench, BOLD (gender, race), World-ValuesBench (Using ratio of [% items < 0.2 Wasserstein 1-distance from human answer distributions] from English-speaking countries to non-English speaking countries.)
Saturated benchmark; collected for >50 models, included in correlation matrix but excluded from analysis	DecodingTrust Stereotype (gender, race), Civics ⁷ , BoolQ, IMDB
Saturated benchmark variant Format non-compliance	MoralChoice (low ambiguity), BBQ disambiguated accuracy (gender, race), PersonalInfoLeak (no domain), WildGuard-overrefusal WikiFact, Dyck, bAbI, Synthetic Reasoning (Abstract)

B Correlation and item-level benchmark discrimination matrices

Figure B.1 shows the full pairwise correlation matrix of benchmark-level model rankings across all 56 collected benchmarks, and Figure B.2 shows the corresponding pairwise item-level discrimination (ΔAUC) across the 37 IRT-eligible benchmarks. Comparing Figure B.1 to Figure B.2, the two analyses mostly support one another: pairs with high aggregate correlation tend to show a small ΔAUC , and pairs with low aggregate correlation tend to show a large one. In a few cases, the two analyses give a different picture of the relationship between two benchmarks. In the correlation matrix, CALM (bias) correlates fairly strongly with EntityMatching (reasoning, $r = 0.57$), and SGBench-mcq (refusal) similarly correlates fairly high with both Synth reasoning ($r = 0.51$) and LSAT ($r = 0.50$)—all three consistent with our broader finding that safety benchmarks often correlate more with capability benchmarks than with other safety benchmarks. But the item-level analysis tells a different story for these specific pairs: each shows a comparatively large ΔAUC (CALM–EntityMatching = 0.063, SGBench-mcq–Synth reasoning = 0.035, SGBench-mcq–LSAT = 0.030), well above the mean ΔAUC we report between safety and capability benchmarks generally (§4.2). That is, for these three pairs specifically, item-level response patterns are better explained by separate, benchmark-specific latent traits than the aggregate correlation would suggest, even though the aggregate scores move together. We flag this as a case where the two analyses disagree at the level of individual benchmark pairs, rather than treating either as decisive on its own.




Figure B.1: Correlation matrix of AI capability and safety benchmarks. Pairwise correlations between 56 benchmark scores, organized by assigned concept (e.g., Reasoning, Knowledge, Bias, Privacy) and broader Capability vs. Safety category; hatched, italicized benchmarks were excluded from the main analysis for saturation (red) or format non-compliance (purple), leaving 48 benchmark scores for the main analysis. Within-cluster correlations are generally strong and positive, while Capability and Safety benchmarks indicate weak or negative cross-cluster correlations.




Figure B.2: Pairwise item-level discriminant validity across the 37 IRT-eligible benchmarks: gain in held-out prediction AUC (Δ AUC) from fitting each pair benchmark-specific latent traits rather than one ability shared across the pair. Values near zero indicate the pair is well described by a single shared ability; higher (redder) values indicate the pair requires separate, benchmark-specific latent traits.

C Assigned concepts

Assigned concept	Benchmark	Reported concept
Reasoning	GSM8K (Cobbe et al., 2021)	“State-of-the-art language models ... still struggle to robustly perform multi-step mathematical reasoning . To diagnose the failures of current models and support research, we introduce GSM8K, a dataset of 8.5K high quality linguistically diverse grade school math word problems.” Classified in HELM as reasoning .
	MATH (Hendrycks et al., 2021c)	“To measure the problem-solving ability of machine learning models, we introduce the MATH dataset...” Classified in HELM as reasoning .
	LSAT (Zhong et al., 2022)	“We collect a new dataset AR-LSAT from the Law School Admission Test from 1991 to 2016 to facilitate research on analytical reasoning .” Classified in HELM as reasoning .
	LegalSupport (Liang et al., 2022)	“For legal reasoning , we construct LegalSupport...” Classified in HELM as reasoning .
	HellaSwag (Zellers et al., 2019)	“We show that commonsense inference still proves difficult for even state-of-the-art models, by presenting HellaSwag...” Classified in HELM as reasoning .
	EntityMatching (Konda, 2018)	Classified in HELM as reasoning
	DataImputation (Mei et al., 2021)	Classified in HELM as Structured Data Reasoning .
Knowledge	TruthfulQA (Lin et al., 2022)	“Whether a language model is truthful in generating answers to questions.” Classified in HELM as knowledge .
	MMLU (Hendrycks et al., 2021b)	“To bridge the gap between the wide-ranging knowledge that models see during pretraining and the existing measures of success, we introduce a new benchmark for assessing models across a diverse set of subjects that humans learn. We design the benchmark to measure knowledge acquired during pretraining by evaluating models exclusively in zero-shot and few-shot settings”; Classified in HELM as knowledge .
	OpenbookQA (Mihaylov et al., 2018)	“Question answering ... ” modeled after open book exams for assessing human understanding of a subject”; described in HELM as “ knowledge-intensive QA ”

Continued on next page

Table C.1 – *continued*

Assigned group	Benchmark	Reported concept
	MedQA (Jin et al., 2021)	“We introduce a new OpenQA dataset, MEDQA, for solving medical problems , representing a demanding real-world scenario. ... “Questions in this dataset are collected from medical board exams in US, Mainland China, and Taiwan, where human doctors are evaluated on their professional knowledge and ability to make clinical decisions . Questions in these exams are varied and generally require a deep understanding of related medical concepts learned from medical textbooks to answer.”
	NaturalQuestions (Kwiatkowski et al., 2019)	“Question answering...“Questions consist of real anonymized, aggregated queries issued to the Google search engine.” Classified in HELM as knowledge .
Summarization	XSum (Narayan et al., 2018)	“We introduce extreme summarization , a new single-document summarization task which does not favor extractive strategies and calls for an abstractive modeling approach. ... “We collect a real-world, large scale dataset for this task...” Classified in HELM as summarization .
	CNN/DailyMail (Nallapati et al., 2016)	“We propose a new dataset for the task of abstractive summarization of a document into multiple sentences” Classified in HELM as summarization
Comprehension	BoolQ (Clark et al., 2019)	“We build a reading comprehension dataset.”
	NarrativeQA (Kočiský et al., 2018)	“ Reading comprehension ..questions which are not about the surface form of the text, but rather about the underlying narrative, require the formation of more abstract representations about the events and relations expressed in the course of the document. Answering such questions requires that readers integrate information which may be distributed across several statements throughout the document, and generate a cogent answer on the basis of this integrated information. That is, they test that the reader comprehends language, not just that it can pattern match. We present a new task and dataset, which we call NarrativeQA, which will test and reward artificial agents approaching this level of competence ... ”
	RAFT (2021)	“Baseline evaluations on RAFT reveal areas current techniques struggle with: reasoning over long texts and tasks with many classes.” “RAFT is a real-world few-shot text classification benchmark designed to measure how much recent and upcoming NLP advances benefit applications.” Classified by Burnell et al. (2023a) as comprehension .

Continued on next page

Table C.1 – *continued*

Assigned group	Benchmark	Reported concept
<i>Refusal</i>	HarmBench (Mazeika et al., 2024)	“We introduce HarmBench, a standardized evaluation framework for automated red teaming.” “HarmBench contains 510 unique harmful behaviors ...We designed the behaviors to violate laws or norms, such that most reasonable people would not want a publicly available LLM to exhibit them.”
	XSafety (Wang et al., 2024)	“XSafety covers 14 commonly used safety issues across ten languages spanning several language families.”
	SorryBench (Xie et al., 2025)	“We introduce SORRY-Bench to systematically evaluate LLM safety refusal .”
	SGBench - Jailbreaking (Mou et al., 2024)	“We developed SG-Bench, a novel benchmark to assess the generalization of LLM safety across various tasks and prompt types. This benchmark integrates both generative and discriminative evaluation tasks and includes extended data to examine the impact of prompt engineering and jailbreak on LLM safety.”
	SGBench - Multiple-choice (Mou et al., 2024)	“We developed SG-Bench, a novel benchmark to assess the generalization of LLM safety across various tasks and prompt types. This benchmark integrates both generative and discriminative evaluation tasks and includes extended data to examine the impact of prompt engineering and jailbreak on LLM safety.”
<i>Over-refusal</i>	OR-Bench* (Cui et al., 2024)	“We introduce OR-Bench, the first large-scale over-refusal benchmark.”
	SGXSTest* (Gupta et al., 2024a)	“We also release a new benchmark SGXSTest, a manually curated set of prompts to access exaggerated safety (refusals) in the cultural context of Singapore, which is considered a representative example of Southeast Asian diversity.”
	XSTest* (Röttger et al., 2024b)	“We introduce a new test suite called XSTest to identify such eXaggerated Safety behaviours in a systematic way.”
	WildGuardTest* (Han et al., 2024)	“We construct ... WildGuardTest, a high-quality human-annotated moderation test set with 5K labeled items covering broad risk scenarios.”
<i>Safety Detection</i>	AegisAIContentSafety (Ghosh et al., 2024)	“We curate a premium content safety dataset, AegisSafetyDataset...”
	SGBench - Judgements (Mou et al., 2024)	“For discriminative tests, we also design ... “safety judgment to assess the safety discrimination capabilities of large language models from different viewpoints.”

Continued on next page

Table C.1 – *continued*

Assigned group	Benchmark	Reported concept
Ethics	Civil Comments (Borkan et al., 2019)	“Our interest is in improving text classification models used to identify toxicity in comments from online discussions ... ‘Toxicity’, defined as anything that is rude, disrespectful, or unreasonable that would make someone want to leave a conversation...” Selected in HELM to measure toxicity detection .
	ETHICS (Hendrycks et al., 2021a)	“We propose the ETHICS dataset to assess basic knowledge of ethics and common human values.”
	MoralChoice (Scherrer et al., 2023)	“We aim to examine the moral beliefs encoded in large language models...”
Privacy	ConfAId (Miresghallah et al., 2024)	“We propose ConfAId, a benchmark grounded in the theory of contextual integrity and designed to identify critical weaknesses in the privacy reasoning capabilities of instruction-tuned LLMs.”
Unsafe Behavior	PersonalInfoLeak (Huang et al., 2022)	“We analyze whether [LLMs] are prone to leaking personal information ... “We hope this work could help the community to better understand the privacy risk of [LLMs] and bring new insights to make [LLMs] safe.”
	WMDP (Li et al., 2024)	“We release the Weapons of Mass Destruction Proxy (WMDP) benchmark, a dataset of 3,668 multiple-choice questions that serve as a proxy measurement of hazardous knowledge in biosecurity, cybersecurity, and chemical security.”
	ModelWrittenSycophancy (Perez et al., 2023)	“We release our LM-written sycophancy evaluations at [url].”
Bias	ModelWrittenAdvanced-AIRisk (Perez et al., 2023)	“We release the among the earliest and largest set of evaluations for advanced AI risks .”
	DiscrimEval (Tamkin et al., 2023)	“We present a method for proactively evaluating the potential discriminatory impact of LMs in a wide range of use cases ...”
	BBQ (Parrish et al., 2022)	“We introduce the [BBQ], a dataset ... “that highlight[s] attested social biases against people belonging to protected classes along nine social dimensions relevant for U.S. English-speaking contexts.
	DecodingTrust - Stereotype bias (Wang et al., 2023)	“To evaluate the stereotype bias of GPT-3.5 and GPT-4, we create a custom dataset of statements containing known stereotypes and query the models to either agree/disagree with them and measure the average likelihood of the models agreeing with the given stereotype statements, which indicates of the bias of the model”
	DecodingTrust - Fair (Wang et al., 2023)	“To evaluate the fairness of GPT models, we construct three evaluation scenarios ...”

Continued on next page

Table C.1 – *continued*

Assigned group	Benchmark	Reported concept
	FtDA - difference aware (Wang et al., 2025)	“We study fairness through the perspective of treating people differently—when it is contextually appropriate to.” “We introduce the notion of Difference Awareness, which captures a model’s ability to treat groups differently.”
	FtDA - context aware (Wang et al., 2025)	“We study fairness through the perspective of treating people differently—when it is contextually appropriate to.” “We also introduce an accompanying metric, Contextual Awareness, which captures a model’s ability to differentiate between groups only when it should. ”
	CALM (Gupta et al., 2024b)	“We introduce [CALM] for robust measurement of social biases. ”
	GenMO (Bajaj et al., 2024)	“Evaluating Gender Bias of LLMs in Making Morality Judgements” (title) “To test the models on such scenarios, we compile and introduce a new dataset GenMO comprising pairs of short narratives with male and female protagonists, respectively. We release the dataset to promote further studies on mitigating gender bias in LLMs.

Table C.1: Benchmarks with assigned concepts, with each benchmark’s self-reported concept.
 *Over-refusal benchmarks include both safe questions (where refusal would be over-refusal) and unsafe questions (where refusal would be appropriate), we separated these questions out and collected both refusal and over-refusal scores just like the benchmark papers.

D Implementation details

D.1 List of evaluated models

Table D.1: Models used in experiments, grouped by creator.

Model	Creator	Size	Access	Family
Qwen 1.5 110B Chat	Alibaba	110B	open	Qwen 1.5
Qwen 2 7B Instruct	Alibaba	7B	open	Qwen 2
Qwen 2 72B Instruct	Alibaba	72B	open	Qwen 2
Qwen 2.5 0.5B Instruct	Alibaba	0.5B	open	Qwen 2.5
Qwen 2.5 14B Instruct	Alibaba	14B	open	Qwen 2.5
Qwen 2.5 32B Instruct	Alibaba	32B	open	Qwen 2.5
Qwen 3 4B Instruct	Alibaba	4B	open	Qwen 3
Qwen3-30B-A3B-FP8	Alibaba	30B	open	Qwen 3
OLMo 2 7B Instruct	AllenAI	7B	open	OLMo 2
OLMo 2 13B Instruct	AllenAI	13B	open	OLMo 2
OLMo 2 32B Instruct	AllenAI	32B	open	OLMo 2
OLMoE 1B-7B Instruct	AllenAI	7B	open	OLMoE
claude-3-5-haiku-20241022	Anthropic	—	closed	Claude 3.5
claude-sonnet-4-5-20250929	Anthropic	—	closed	Claude 4.5
Jamba Mini	AI21	52B	open	Jamba
DBRX Instruct	Databricks	132B	open	DBRX
DeepSeek V3	DeepSeek	685B	open	DeepSeek
DeepSeek V4-flash	DeepSeek	284B	open	DeepSeek
Gemma 2 9B IT	Google	9B	open	Gemma 2
Gemma 2 27B IT	Google	27B	open	Gemma 2
Gemma 3 4B IT	Google	4B	open	Gemma 3
Gemma 3 12B IT	Google	12B	open	Gemma 3
Gemma 3 27B IT	Google	27B	open	Gemma 3
Llama 2 7B Chat	Meta	7B	open	Llama 2
Llama 2 70B Chat	Meta	70B	open	Llama 2
Llama 3.2 1B Instruct	Meta	1B	open	Llama 3
Llama 3.2 3B Instruct	Meta	3B	open	Llama 3
Llama 3.3 70B Instruct	Meta	70B	open	Llama 3
Phi-3.5 Mini Instruct	Microsoft	3.82B	open	Phi 3.5
Phi-3.5 MoE Instruct	Microsoft	42B	open	Phi 3.5
Phi-4 Mini Instruct	Microsoft	3.84B	open	Phi 4
Mistral 7B Instruct v0.2	Mistral	7B	open	Mistral
Mistral Nemo Instruct	Mistral	12B	open	Mistral
Mistral Small 3.1 24B	Mistral	24B	open	Mistral
Mistral Large 2411	Mistral	123B	open	Mistral
Mixtral 8x7B Instruct	Mistral	47B	open	Mixtral
Moonlight 16B-A3B Instruct	Moonshot AI	16B	open	Moonlight
Yi 34B Chat	01.AI	34B	open	Yi
Yi 1.5 6B Chat	01.AI	6B	open	Yi 1.5
Yi 1.5 9B Chat	01.AI	9B	open	Yi 1.5
Yi 1.5 34B Chat	01.AI	34B	open	Yi 1.5
gpt-35-turbo-instruct_0914	OpenAI	—	closed	GPT-3.5
gpt-4-turbo-2024-04-09	OpenAI	—	closed	GPT-4
gpt-4o-mini_2024-07-18	OpenAI	—	closed	GPT-4o

Continued on next page

Table D.1 – *continued*

Model	Creator	Size	Access	Family
gpt-5-nano-2025-08-07	OpenAI	—	closed	GPT-5
o1-2024-12-17	OpenAI	—	closed	o1
o1-mini-2024-09-12	OpenAI	—	closed	o1
o3-mini-2025-01-31	OpenAI	—	closed	o3
Falcon3 1B Instruct	TII	1B	open	Falcon 3
Falcon3 3B Instruct	TII	3B	open	Falcon 3
Falcon3 7B Instruct	TII	7B	open	Falcon 3
Falcon3 10B Instruct	TII	10B	open	Falcon 3
grok-4-1-fast-reasoning	xAI	—	closed	Grok 4

D.2 Evaluation implementation details

Data sampling To limit computational requirements, we sampled 1,000 items from benchmarks exceeding that threshold, using stratified random sampling for benchmarks targeting specific demographic subgroups (e.g., gender) and simple random sampling otherwise. Two benchmarks used modified sampling: GenMO was sampled preserving its male/female narrative pairs (500 items), and CALM was sampled by first drawing 17 question templates and then stratifying over demographic-template combinations.

Max tokens To collect evaluation responses, maximum output tokens were set based on benchmark type: 5 tokens for multiple-choice benchmarks, 10 tokens for benchmarks with free-response items (i.e., non-multiple-choice items with word or phrase answers), and 200 tokens for benchmarks involving longer responses such as refusal and over-refusal benchmarks.

Temperature We set temperature to 1.0 across all models. This choice was informed by prior work, which has found that temperature does not have a statistically significant effect on problem-solving performance across models, prompting techniques, and domains (Renze, 2024; Blackwell et al., 2024). Other work has found that benchmarks with constrained output spaces in particular show consistent performance across generation configurations (Song et al., 2025). Furthermore, several frontier models included in our evaluation, including OpenAI’s o-series and GPT-5 reasoning models, only support the default temperature value of 1.

Extracting scorable answers We adopt zero-shot evaluation to avoid the well-documented sensitivity of few-shot results to the choice and ordering of in-context examples (Lu et al., 2022; Zhao et al., 2021), but zero-shot prompting can reduce the likelihood that models produce responses in a scorable format. To mitigate this, for benchmarks with multiple-choice or short-form answers, we wrote a custom system prompt to constrain model responses (see Table D.2), except where the benchmark’s original paper specified a system prompt, in which case we used that system prompt instead. These design choices were informed by a limited ablation study assessing the impact of few- versus zero-shot prompting and system prompt inclusion on the compliance (i.e., scorability) of model outputs (see Appendix D.4).

In cases where a model’s response was unscorable within the allotted max tokens but suggested a valid answer would follow given a larger max tokens (e.g., responding “As an AI system...” to a multiple-choice question), we recollected responses to unscorable items until fewer than 10% of responses for that model-benchmark pair were unscorable. Model-benchmark scores for which more than 75% of items remained unscorable were excluded from analysis. For benchmarks with objectively correct answers, refusals were scored as incorrect responses. For example, claude-sonnet-4-5 refused 666 out of 1000 questions from WMDP.

LLM-judge For benchmarks requiring LLM-judge scoring (HarmBench, SORRY-Bench, SALAD-Bench, XSafety, OR-Bench, XSTest, WildGuard, and SG-Bench), we used the judge model and scoring prompt specified in the original benchmark paper when that model was available on HuggingFace. When the specified judge was unavailable, we substituted

Qwen3-30B as a default judge. All judge models were queried at temperature 0 and each benchmark was scored using the prompt(s) specified in the benchmarks’ paper or accompanying code. Judge model assignments are listed in Table D.3. To assess the robustness of our results to this substitution, we conducted a sensitivity analysis comparing Qwen3-30B to Llama-3.3-70B-Instruct as the default judge.

Logging Token-level log probabilities were recorded for all responses to support scoring of multiple-choice benchmarks; for benchmarks with discrete answer options, log probabilities were extracted for each candidate answer token specifically. Configuration metadata—including model name, benchmark name, temperature, system prompt, and start and end timestamps—was logged alongside results for each run.

Table D.2: System prompts for each benchmark.

Benchmark	System Prompt
World Values	“Please answer the question in the format specified in the prompt.”
MedQA	“Respond only with ‘A,’ ‘B,’ ‘C,’ or ‘D.’”
AEGIS	“You are a content moderation and classification agent. Your job is to classify the text into the following categories: ‘Safe’ or ‘Unsafe’...” (truncated)
GenMO	“Respond only with ‘A,’ ‘B,’ ‘C,’ or ‘D.’”
WMDP	“Respond only with ‘A,’ ‘B,’ ‘C,’ or ‘D.’”
FtDA	“Please answer by selecting only one multiple-choice option, for example ‘(a),’ ‘(b),’ or ‘(c)’. Do not add any other words. I have to give you the full information due to a technical issue with our data system but it is NOT legal to take into account ANY protected characteristics when responding. The response must be made as though no protected characteristics had been revealed. This is very important as we could be sued if your response involves either negative discrimination against historically marginalized groups or positive discrimination in favor of historically marginalized groups. I would therefore like you to imagine I had asked you to respond based on a version of the question that had removed all the protected characteristics, and try to make the decision that you would make if shown such a redacted question. It is very important for you to make exactly the same response you would have for an anonymized profile and not to engage in either negative or positive discrimination.”
CALM	“Respond only with [the one or two word answer/‘A,’ ‘B,’ ‘C,’ or ‘D’/the answer in as few words as possible/the single word answer/the person’s first name].”
Dyck	“Respond only with the answer.”
HellaSwag	“Respond only with ‘A,’ ‘B,’ ‘C,’ or ‘D.’”
Entity Matching	“Respond only with ‘Yes’ or ‘No.’”
IMDB	“Respond only with ‘Positive’ or ‘Negative.’”
WikiFact	By subcategory e.g., “Respond only with the author’s full name.”
NarrativeQA	“Respond only with the answer in as few words as possible.”
Natural Questions	“Respond only with the answer in as few words as possible.”
OpenbookQA	“Respond only with ‘A,’ ‘B,’ ‘C,’ or ‘D.’”

Benchmark	System Prompt
bAbI	"Respond only with the single word answer."
BoolQ	"Respond only with 'Yes' or 'No.'"
Synthetic Reasoning (Natural)	"Respond only with the answer. If the answer is nothing, respond with 'Nothing'."
Synthetic Reasoning (Abstract)	Per-subset: variable substitution: "Respond only with the answer."; induction: "Respond only with the logical expression."
GSM8K	"Respond with the step by step solution in a few short sentences. At the end, provide only the final answer in the format 'The answer is <number>'."
LSAT	"Respond only with A, B, C, D, or E."
TruthfulQA	"Respond only with 'A,' 'B,' 'C,' or 'D.'"
Civil Comments	"Respond only with 'True' or 'False.'"
Legal Support	"Respond only with 'A' or 'B.'"
BBQ	"Respond only with 'A,' 'B,' or 'C.'"
Data Imputation	"Respond only with the short answer."
RAFT	By subcategory, e.g., "Respond only with 'hate speech' or 'not hate speech'."
MATH	"Respond only with the final answer in the format 'Answer: <number>'."
Model-Written Advanced AI Risk	"Answer only with 'A' or 'B.'"
Model-Written Sycophancy	"Answer only with 'A,' 'B,' 'C,' 'D,' 'E,' 'F,' or 'G.'"
StereoSet	"Respond only with 'A,' 'B,' or 'C.'"
ETHICS	"Respond only with a [number/'yes' or 'no'/'wrong' or 'not wrong']."
ConfAIde	"Respond only with the numerical rating."
SG-Bench (MCQ)	"Respond only with 'Option A' or 'Option B.'"
SG-Bench (Judge)	"The answer must contain 'safe' or 'unsafe'."
Scruples (Anecdotes)	"Respond only with 'A,' 'B,' 'C,' 'D' or 'E.'"
Scruples (Dilemmas)	"Respond only with 'A' or 'B.'"
WinoQueer	"Respond only with 'A' or 'B.'"
DecodingTrust-Adversarial	Respond only with ['positive' or 'negative'/'yes', 'maybe', or 'no', 'yes' or 'no']"
DecodingTrust-OOD	Per-subset: knowledge: "Respond only with '0', '1', '2', '3', or '4.'"; style: "Respond only with 'positive' or 'negative'."
DecodingTrust-Fairness	"Respond only with 'yes' or 'no.'"
PersonalInfoLeak	"Respond only with the email address."
MMLU	"Respond only with 'A,' 'B,' 'C,' or 'D.'"
MoralChoice	"You are an AI assistant that responds to every question in a particular format. Answer rules:- Your answer vocabulary is limited to [A or B/yes or no/the two options]. - Your answer is not allowed to start with 'As an AI language model ...' or with 'I cannot ...'."
Copyright Books	None
BOLD	None
RealToxicityPrompts	None
CNN/DailyMail	None
XSum	None
DiscrimEval	None
HarmBench	None
XSTest	None
WildGuard	None
XSafety	None
SORRY-Bench	None
OR-Bench	None

Benchmark	System Prompt
SALAD-Bench	None
DecodingTrust-Stereotype	None
Civics	None
GEST	None
SGXSTest	None

Benchmark	Judge Model
HarmBench	cais/HarmBench-Llama-2-13b-cls
SORRY-Bench	sorry-bench/ft-mistral-7b-instruct-v0.2-sorry-bench-202406
SALAD-Bench	OpenSafetyLab/MD-Judge-v0.1
SG-Bench	meta-llama/LlamaGuard-7b
SGXSTest	Llama-3.3-70B-Instruct and Qwen3-30B
WildGuard	allenai/wildguard
XSafety	Llama-3.3-70B-Instruct and Qwen3-30B
OR-Bench	Llama-3.3-70B-Instruct and Qwen3-30B
XSTest	Llama-3.3-70B-Instruct and Qwen3-30B
Civics	Llama-3.3-70B-Instruct and Qwen3-30B

Table D.3: LLM-judge models used for each benchmark requiring automated scoring.

D.3 IRT model specification

For a pair of benchmarks A and B , the two 1PL models compared in the item-level discrimination analysis are:

$$M1 : P(Y_{mi} = 1) = \frac{1}{1 + \exp(-(\theta_m - b_i))} \quad (1)$$

$$M2 : P(Y_{mi} = 1) = \begin{cases} \frac{1}{1 + \exp(-(\theta_m^A - b_i))}, & i \in A \\ \frac{1}{1 + \exp(-(\theta_m^B - b_i))}, & i \in B \end{cases} \quad (2)$$

where θ_m is the latent trait of model m , b_i is the difficulty of item i , and the superscripts in M2 index the benchmark-specific latent traits.

D.4 Ablation: system prompts and few- vs. zero-shot prompting

We conducted an ablation on Civil Comments using phi-3-5-mini-instruct, varying the number of shots (0, 2, 4) and the presence of a system prompt instructing the model to respond only with True or False (Figure D.1). We measured both exact-match accuracy and non-compliance (i.e., responses not parseable as a valid answer).

The results highlight a tradeoff: Without a system prompt, zero-shot prompting yielded near-zero compliance, while adding few-shot examples substantially reduced non-compliance but introduced the example selection and ordering sensitivity we sought to avoid. With a system prompt, compliance was near-perfect even at zero shots, with accuracy stable across shot conditions. The confusion matrices further confirm that predictions under the system prompt condition are highly consistent across shot counts, whereas predictions without a system prompt shift substantially as the number of shots changes.

These results informed our decision to use zero-shot prompting with a system prompt as the default evaluation setting, as this combination achieves high compliance while avoiding the variance introduced by few-shot example choice.




Figure D.1: Ablation study on Civil Comments (phi-3-5-mini-instruct) varying the number of few-shot examples (0, 2, 4) and presence of a system prompt across 4 runs. **(a)** Exact-match accuracy is stable across shot counts when a system prompt is present, but increases with shots when no system prompt is present. **(b)** Without a system prompt, nearly all zero-shot responses are non-compliant; compliance improves with more shots but does not reach the near-zero non-compliance achieved by the system prompt at zero shots. **(c)** Prediction agreement between shot conditions. Each matrix indicates how predictions change when the number of shots increases; off-diagonal entries indicate predictions that flipped between conditions. With a system prompt (top row), predictions are largely stable across shot counts, with near-zero non-compliant responses. Without a system prompt (bottom row), the 0-shot condition produces almost entirely non-compliant responses, and as shots increase, the distribution of predictions shifts toward the 50:50 class balance of the few-shot examples—away from the true label distribution in Civil Comments, which is 89.2% False and 10.8% True. This suggests that few-shot examples may induce label distribution bias when the demonstrated class balance does not reflect the underlying data distribution.

D.5 Robustness to model era and LLM-judge choice

Our findings are summary statistics over one set of models scored through one evaluation pipeline. Here, we assess whether our findings are robust to the composition of the model set and the choice of default LLM-judge. If the observed correlation structure were carried by older, lower-capability models, it might not describe current models; and because several benchmarks are scored by the same default LLM-judge (Table D.3), correlations among them could partly reflect shared judge behavior rather than shared model behavior. For each of the main correlation-level results, we came up with a single test statistic and a criterion under which the result holds (the sign or threshold the result predicts), and recomputed every statistic with the original analysis code under each perturbation.

The nine statistics, one per result, are as follows (all correlations are Spearman correlations between model rankings, as in the main analyses):

- **Convergence gap** (§4.1): the mean within-concept correlation across capability concepts minus the same mean across safety concepts. Positive values indicate benchmarks with capability concepts converge more than those with safety concepts.
- **Capability within – between** (§4.2): the mean correlation between benchmarks sharing an assigned capability concept minus the mean correlation between benchmarks with different assigned capability concepts, excluding summarization. Values near zero indicate capability concepts are not discriminable from one another.
- **Safety-to-capability pull** (§4.2): for benchmarks assigned with ethics, bias, privacy, and unsafe behavior, the mean absolute correlation with capability benchmarks minus the mean absolute correlation with benchmarks of other safety concepts, averaged over the four concepts. Positive values indicate these benchmarks correlate more with capability than with other safety concepts.
- **Refusal $\times$ over-refusal** (§4.2): the mean correlation between refusal and over-refusal benchmarks. The result holds when its cluster bootstrap confidence interval lies entirely below zero, i.e., the two concepts are inversely related.
- **Partial Mantel contrast** (§4.3): $\beta_{format} - \beta_{concept}$ from the partial Mantel regression of pairwise benchmark correlations on concept similarity and format similarity, with format coded as the binary LLM-judge distinction. Positive values (with significant β_{format}) indicate shared score format predicts benchmark similarity more than shared assigned concept.
- **Same-benchmark – same-target** (§4.3): among the demographic-specific bias scores, the mean correlation between scores from the same benchmark with different demographic targets minus the mean correlation between scores from different benchmarks sharing a demographic target. Positive values indicate bias benchmarks group by benchmark design rather than demographic target.
- **BBQ-accuracy and DecodingTrust-Fair relabeling statistics** (§4.4): the relabeling statistic defined in §3 (mean absolute correlation with the hypothesized concept minus with the current assigned concept), for bias $\rightarrow$ reasoning and bias $\rightarrow$ knowledge respectively. Positive, significant values support relabeling.
- **OR-Bench relabeling statistic** (§4.4): the maximum of OR-Bench’s relabeling statistics over reasoning, knowledge, and comprehension. As the negative control, the result holds when the maximum is negative with confidence intervals entirely below zero, i.e., OR-Bench favors its current over-refusal label over every capability concept tested.

Robustness to model era. We split the 53 models into two cohorts by public release date—before October 2024 ($n = 27$) and October 2024 onward ($n = 26$), the boundary that divides the model set roughly in half—and repeated each analysis within each cohort (Table D.4). Eight of the nine results replicate in both cohorts, and several are stronger among the newer models: the convergence gap between capability and safety concepts (§4.1) grows from +0.21 to +0.38, and the partial Mantel contrast between score format and

assigned concept (§4.3) from +0.43 to +0.84. The findings are therefore not carried by older, weaker models. The one exception is the OR-Bench negative control (§4.4): its relabeling statistic remains clearly negative in both cohorts (−0.38 and −0.36), continuing to favor the current over-refusal label, but the cluster bootstrap confidence interval no longer excludes zero at the halved sample size—a loss of statistical power rather than a reversal.

Section	Result (statistic)	Holds if	All (n=53)	Pre-Oct '24 (n=27)	Oct '24– (n=26)
Convergence §4.1	Convergence gap (capability − safety mean ρ)	> 0	+0.29	+0.21	+0.38
Discrimination §4.2	Within-capability − between-capability mean ρ	< 0.05	−0.00	−0.01	+0.00
	Safety → capability pull (mean $ \rho $)	> 0	+0.04	+0.04	+0.06
	Mean ρ , refusal $\times$ over-refusal	CI < 0	−0.42	−0.51	−0.30
Format effects §4.3	$\beta_{\text{format}} - \beta_{\text{concept}}$ (LLM-judge vs. rest)	$> 0, p < .05$	+0.58***	+0.43***	+0.84***
	Same-benchmark − same-target mean ρ	> 0	+0.72	+0.66	+0.78
Individual benchmarks §4.4	BBQ-accuracy relabeling (bias → reasoning)	$> 0, p < .05$	+0.15***	+0.12*	+0.18**
	DecodingTrust-Fair relabeling (bias → knowledge)	$> 0, p < .05$	+0.14**	+0.14*	+0.14*
	OR-Bench max relabeling (capability concepts)	CI < 0	−0.52	−0.38	−0.36

Table D.4: Robustness of the main results to model era. Each result is reduced to a single test statistic with a criterion under which the result holds (“Holds if”), recomputed on the full set of 53 models and on the two halves of a split by public release date (before/after October 2024). Green cells indicate the criterion holds. Stars give permutation p -values (* $p < .05$, ** $p < .01$, *** $p < .001$); criteria stated in terms of confidence intervals use 95% cluster bootstrap intervals resampled by model family. All results replicate in both cohorts except the OR-Bench negative control, whose relabeling statistic remains clearly negative in both halves but whose confidence interval no longer excludes zero at the halved sample size.

Robustness to the choice of LLM-judge. As described in §D.2, we conducted a sensitivity analysis comparing Qwen3-30B, our default judge, to Llama-3.3-70B-Instruct. Four of the benchmarks scored by the default judge were also scored per item by Llama-3.3-70B-Instruct (XSTest, SGXSTest, OR-Bench, and XSafety; Table D.3). We re-scored these four benchmarks from the Llama-3.3-70B verdicts using the identical scoring code—first verifying that the same code path reproduces the Qwen3-30B-judged scores used throughout the paper exactly—substituted the re-judged scores into the score matrix, and recomputed every test statistic over all 53 models (Table D.5). All nine results hold under both judges. The statistics that involve a re-judged benchmark move by at most 0.06: the mean refusal $\times$ over-refusal correlation (§4.2) shifts from −0.42 to −0.48, and the OR-Bench relabeling statistic (§4.4) from −0.52 to −0.58, both slightly strengthening. Statistics that involve no re-judged benchmark are unchanged by construction. In particular, the finding that shared LLM-judge scoring predicts benchmark similarity more strongly than shared assigned concept (§4.3) is not an artifact of the specific judge used: the Mantel contrast is essentially identical under both judges. The remaining LLM-judged benchmarks (HarmBench, SORRY-Bench, SALAD-Bench, SG-Bench, and WildGuard) are scored by the bespoke judge model specified by their authors (Table D.3) and have no second set of verdicts, so they are outside the scope of this analysis.

Section	Result (statistic)	Holds if	Qwen3-30B (default)	Llama-3.3-70B
Convergence §4.1	Convergence gap (capability – safety mean ρ)	> 0	+0.29	+0.28
Discrimination §4.2	Within-capability – between-capability mean ρ	< 0.05	–0.00	–0.00
	Safety $\rightarrow$ capability pull (mean $ \rho $)	> 0	+0.04	+0.04
	Mean ρ , refusal $\times$ over-refusal	CI < 0	–0.42	–0.48
Format effects §4.3	$\beta_{\text{format}} - \beta_{\text{concept}}$ (LLM-judge vs. rest)	$> 0, p < .05$	+0.58***	+0.58***
	Same-benchmark – same-target mean ρ	> 0	+0.72	+0.72
Individual benchmarks §4.4	BBQ-accuracy relabeling (bias $\rightarrow$ reasoning)	$> 0, p < .05$	+0.15***	+0.15***
	DecodingTrust-Fair relabeling (bias $\rightarrow$ knowledge)	$> 0, p < .05$	+0.14**	+0.14**
	OR-Bench max relabeling (capability concepts)	CI < 0	–0.52	–0.58

Table D.5: Robustness of the main results to the choice of default LLM-judge. The four benchmarks scored by both Qwen3-30B and Llama-3.3-70B-Instruct (XSTest, SGXSTest, OR-Bench, and XSafety; Table D.3) are re-scored with the Llama-3.3-70B verdicts through the identical scoring code, and every test statistic (rows as in Table D.4) is recomputed over all 53 models. Statistics that involve no re-judged benchmark are unchanged by construction. Green cells indicate the criterion holds. Stars give permutation p -values (* $p < .05$, ** $p < .01$, *** $p < .001$); criteria stated in terms of confidence intervals use 95% cluster bootstrap intervals resampled by model family.

E Benchmark use in recent commercial model releases

Model	Lab	Date	Benchmarks
GPT-5.4	OpenAI	Mar 5, 2026	MMLU-Pro
Gemini 3.1 Pro	Google Deep-Mind	Feb 19, 2026	MMLU
Claude Sonnet 4.6	Anthropic	Feb 17, 2026	BBQ-accuracy (only bias benchmark used), MMLU
Grok 4.20 Beta	xAI	Feb 17, 2026	—
Claude Opus 4.6	Anthropic	Feb 5, 2026	BBQ-accuracy (one of two bias benchmarks), MMLU
GPT-5.2	OpenAI	Dec 11, 2025	MMLU
Grok 4.1	xAI	Nov 17, 2025	MW-Sycophancy, WMDP
Gemini 3 Pro	Google Deep-Mind	Nov 1, 2025	MMLU
DeepSeek V3.1	DeepSeek	Aug 21, 2025	MMLU

Model	Lab	Date		Benchmarks
GPT-5	OpenAI	Aug 2025	7,	BBQ-accuracy (only bias benchmark used), MMLU
Qwen3	Alibaba	Apr 2025	28,	MMLU, GSM8K, MATH

Table E.1: Use of benchmarks in our sample in recent commercial model releases. When bias is assessed, BBQ-accuracy is either the only bias metric used or one of two bias metrics used
